


M383C Notes

Arun Debray
October 9, 2015

These notes were taken in UT Austin's Math 383c class in Fall 2015, taught by Todd Arbogast. I live-T_EXed them using vim, and as such there may be typos; please send questions, comments, complaints, and corrections to a.debray@math.utexas.edu.



Source: <http://brownsharpie.courtneygibbons.org/?p=78>

Contents

Chapter 1. Normed Linear Spaces and Banach Spaces	4
1.1 General Remarks: 8/26/15	4
1.2 Banach Spaces: 8/28/15	6
1.3 Bounded Linear Operators: 8/31/15	8
1.4 ℓ^p -norms: 9/2/15	11
1.5 ℓ^p and L^p -spaces: 9/4/15	13
1.6 $L^p(\Omega)$ is Banach: 9/9/15	16
1.7 The Hahn-Banach Theorem: 9/11/15	18
1.8 The Hahn-Banach Theorem, II: 9/14/15	20
1.9 Separability: 9/16/15	22
1.10 The Minkowski Functional and the Baire Category Theorem: 9/18/15	23
1.11 The Open Mapping Theorem: 9/21/15	25
1.12 The Uniform Boundedness Principle: 9/23/15	27
1.13 Weak and Weak-* Convergence: 9/25/15	30
1.14 The Banach-Alaoglu Theorem: 9/28/15	32
1.15 The Generalized Heine-Borel Theorems: 9/30/15	34
1.16 The Dual to an Operator: 10/2/15	35
Chapter 2. Inner Product Spaces and Hilbert Spaces	39
2.1 Orthogonality: 10/5/15	39
2.2 Projections: 10/7/15	41
2.3 Orthonormal Bases: 10/9/15	43

Normed Linear Spaces and Banach Spaces

Lecture 1.

General Remarks: 8/26/15

Though the course name is “Methods of Applied Mathematics,” this is a misnomer; the course is really about functional analysis.

The course will use the Canvas website (<http://canvas.utexas.edu/>), and office hours will be after class (modulo lunch), Mondays and Wednesdays from 12:30 to 1:50. Under UT Direct, there’s also a CLIPS page, but that’s less central to the course.

The textbook is a set of course notes; it hasn’t changed much since 2013, so if you have that version, you’ll be fine. They’ll be ready at the copy center by Friday or Monday.

Homework will be due every week, assigned one Friday, and due the next. The first assignment will be due in a little over a week. We’re encouraged to work in groups, but must write up our own individual proofs. Midterms will be weeks 7 and 12, probably, and will be topical; the final, at the end of the semester, will be comprehensive.

In this course, we’ll cover chapters 2 – 5 of the lecture notes. Some elementary topology and Lebesgue integration (the first chapter) will be assumed.

Now, for some math. The professor is an applied mathematician, doing numerical analysis, and more specifically, approximation of differential equations. Functional analysis is useful for that, but also plenty of other fields, even including abstract algebra! Nonetheless, the course will be presented from an applied perspective.

The background is that we’re trying to solve a problem of the form $T(u) = f$. Here, T is a model or differential equation; it’s some kind of operator. f is the data that we’re given, and we want to find the solution u . We use the framework of functional analysis to understand the nature of the functions u and f : their properties and what classes of functions they live in. We also want to know the nature of the operator T . In particular, we’ll focus on cases where T is linear, since anything nonlinear can usually be locally approximated with a linear one. Thus, we should start with the linear case.

The set of all functions is a vector space, of course, so we’re led to study vector spaces. At the undergraduate level, one studies finite-dimensional spaces, but here we’ll use infinite-dimensional ones. Vector spaces also give us the required linearity. But since we also have questions of convergence, we’ll introduce topology, so this course combines algebra and topology.

In this class, \mathbb{F} will denote a field, either \mathbb{R} or \mathbb{C} (a lot of the time, the stuff we’re doing won’t depend on which).

Definition. Let X be a vector space over \mathbb{F} . Then, X is a *normed linear space* (henceforth NLS) if it has a *norm*, a function $\|\cdot\| : X \rightarrow \mathbb{R}^+ = [0, \infty)$ such that for every $x, y \in X$ and $\lambda \in \mathbb{F}$,

- $\|\lambda x\| = |\lambda| \|x\|$,
- $\|x\| = 0$ iff $x = 0$, and
- $\|x + y\| \leq \|x\| + \|y\|$.

The last stipulation is called the *triangle inequality*.

These conditions on the norm mean it’s a measure of size: stretching a vector stretches the norm, the only thing with size 0 is the origin, and the triangle inequality corresponds to the familiar geometric one. It turns out these are the only properties we need to measure size.

Example 1.1.1.

- (1) d -dimensional Euclidean space \mathbb{F}^d comes with a familiar norm: if $x = (x_1, \dots, x_n)$ for $x_j \in \mathbb{F}$, then

$$\|x\| = \sqrt{\sum_{j=1}^d |x_j|^2}.$$

Sometimes, this is simply denoted $|x|$. Thus, whenever we talk about \mathbb{F}^d , we really mean $(\mathbb{F}^d, \|\cdot\|)$, the normed linear space.

- (2) If $a < b$, where $a, b \in [-\infty, \infty]$, let $C([a, b])$ denote the space of continuous functions $f : [a, b] \rightarrow \mathbb{F}$ such that $\sup_{x \in [a, b]} |f(x)|$ is finite.¹ This is indeed a vector space; then, it turns to a normed linear space with the norm

$$\|f\| = \sup_{x \in [a, b]} |f(x)|.$$

Notice that the norm must be finite, which is satisfied here. The first two properties are clearly satisfied, and because the absolute value is a norm on \mathbb{R} , then the triangle equality is also satisfied.

- (3) We can pair $C([a, b])$ with a different norm $\|\cdot\|_{L^1}$, defined by

$$\|f\|_{L^1} = \int_a^b |f(x)| dx.$$

The integral certainly exists, since f is continuous, but it might be infinite; thus, we assume that a and b are finite, so $[a, b]$ is compact, and

$$\int_a^b |f(x)| dx \leq (b-a) \sup_{x \in [a, b]} |f(x)|,$$

so we're bounded. It's also not that hard to show that $\|\cdot\|_{L^1}$ is a norm, as the integral is linear.

We now have two norms on $C([a, b])$; are they “the same?” Though the underlying vector spaces are the same, the measures of size are different, so as normed linear spaces they are not the same.

We can find more examples sitting inside other NLSes.

Proposition 1.1.2. *Let $(X, \|\cdot\|)$ be an NLS and $V \subseteq X$ be a linear subspace. Then, $(V, \|\cdot\|)$ is an NLS.*

It's easy to check that the three requirements are still met.

We can measure size, so since we're in a vector space, we can measure distance. In general, we have a metric. Specifically, if $(X, \|\cdot\|)$ is an NLS, define $d : X \times X \rightarrow \mathbb{R}^+$ by $d(x, y) = \|x - y\|$. Why is this a metric? It has to satisfy the following three properties for all $x, y, z \in X$.

- (1) $d(x, y) = 0$ iff $x = y$.
- (2) $d(x, y) = d(y, x)$.
- (3) $d(x, y) + d(y, z) \geq d(x, z)$.

It's easy to check that the d induced from the norm is indeed a metric; each metric property follows from one of the norm properties.

And now that we can measure distance, we have a topology; specifically a metric topology, the simplest of all topologies. That is, a normed linear space is a metric space. To be specific, define the *ball of radius r about x* , where $r > 0$ and $x \in X$, is

$$B_r(x) = \{y \in X \mid d(x, y) < r\}.$$

This is an open ball, so the distance must be strictly less than r .

The topology is defined by setting $U \subseteq X$ to be *open* if for every $x \in U$, there exists an $r > 0$ such that $B_r(x) \subseteq U$. In other words, an open set doesn't contain its boundary. A set $F \subseteq X$ is *closed* if the complement $F^c = X \setminus F$ is open.

Definition. A subset F of a metric space X is *sequentially closed* if whenever $\{x_n\}_{n=1}^\infty$ is a sequence in F converging to an $x \in X$ (in the sense of the metric, i.e. $d(x_n, x) \rightarrow 0$), then $x \in F$.

¹Recall that the *supremum* of a set is its least upper bound: for example, $\sup(0, 1) = 1$, even though 1 isn't part of the set. This distinguishes the supremum from the maximum.

In a metric space (this is *not* true in general!), F is closed iff F is sequentially closed.

Now, we have algebra (the vector space), the metric (giving us convergence, compactness, etc.), and the norm. How are they related?

Proposition 1.1.3. *In an NLS X , addition, scalar multiplication, and the norm are all continuous functions.*

PROOF. We'll prove this for addition and the norm; scalar multiplication is analogous to addition.

Addition is a function $+: X \times X \rightarrow X$. Let $\{x_n\} \subseteq X$ with $x_n \rightarrow x$ and $\{y_n\} \subseteq X$ with $y_n \rightarrow y$. Continuity is equivalent to $\{x_n + y_n\} \rightarrow x + y$ for all such sequences. That is, I need $d(x_n + y_n, x + y) \rightarrow 0$, but that's equivalent to $\|(x_n + y_n) - (x + y)\| \rightarrow 0$.

Since $x_n \rightarrow x$ and $y_n \rightarrow y$, then $\|x_n - x\| \rightarrow 0$ and $\|y_n - y\| \rightarrow 0$. It looks like we should use the triangle inequality.

$$\begin{aligned} \|(x_n + y_n) - (x + y)\| &= \|(x_n - x) + (y_n - y)\| \\ &\leq \|x_n - x\| + \|y_n - y\| \rightarrow 0. \end{aligned}$$

The norm is a little different. Suppose $x_n \rightarrow x$, which means we need to show that $\|x_n\| \rightarrow \|x\|$. Well,

$$\begin{aligned} \|x\| &= \|x - x_n + x_n\| \\ &\leq \|x - x_n\| + \|x_n\| \\ &\leq 2\|x - x_n\| + \|x\|. \end{aligned}$$

Since we've sandwiched $\|x - x_n\|$, then $\lim \|x_n\| = \|x\|$.² ☒

Lecture 2.

Banach Spaces: 8/28/15

Recall that if $(X, \|\cdot\|)$ is an NLS, we have a metric $d(x, y) = \|x - y\|$ and a topology. More generally, if (X, d) is a metric space, $x_n \rightarrow x$ is the same as $d(x_n, x) \rightarrow 0$. In our case, this means that $\|x_n - x\| \rightarrow 0$.

Definition. A sequence $\{x_n\}_{n=1}^{\infty}$ is a *Cauchy sequence* if $\lim_{n,m \rightarrow \infty} d(x_n, x_m) = 0$.

Here, n and m go to infinity independently, which might be confusing; an alternate way to phrase this is that $\{x_n\}$ is Cauchy if for all $\varepsilon > 0$, there exists an $N = N_\varepsilon > 0$ such that $d(x_n, x_m) \leq \varepsilon$ whenever $m, n \geq N$.

In a Cauchy sequence, the terms get closer and closer together, but do they converge? Consider $(0, \infty)$ and $x_n = 1/n$. This is Cauchy, but would converge to 0, which isn't part of our set; in a sense, it's a "hole" in our set. This is annoying.

Definition.

- A metric space X is *complete* if every Cauchy sequence on X converges in X .
- A complete NLS is called a *Banach space*.

We'll also give some properties of subspaces of NLSes.

Definition. Let X be an NLS. A set $M \subseteq X$ is *bounded* if there exists an $R > 0$ such that $M \subseteq \overline{B_R(0)} = \{x : \|x\| \leq R\}$.

Equivalently, M is bounded if there's a finite R such that $\|x\| \leq R$ for all $x \in M$.

Proposition 1.2.1. *Every Cauchy sequence in an NLS is bounded.*

PROOF. The idea is that all but a finite number of points in a sequence are within distance 1 of each other.

Let $\{x_n\}_{n=1}^{\infty}$ be a Cauchy sequence in an NLS X . By definition (using $\varepsilon = 1$), there's an $N > 0$ such that $\|x_n - x_N\| \leq 1$ for all $n \geq N$. Using the triangle inequality, $\|x_n\| \leq \|x_N\| + 1$ for all $n \geq N$.

Now, let $M = \max\{\|x_1\|, \dots, \|x_{N-1}\|\}$ and $R = \max\{\|x_N\| + 1, M\}$; both of these are finite sets, and therefore have maxima. Thus, $\|x_n\| \leq R$ for all n . ☒

Even if the limit isn't there, the sequence is still bounded, which is nice. Also, notice how we used the norm; boundedness in metric spaces maybe isn't so interesting.

²This was all that the professor said about the proof that the norm is continuous. Here's an alternate proof in case you, like me, didn't get it: since $x_n \rightarrow x$, then for any $n \in \mathbb{N}$, there's an N_n such that if $m \geq N_n$, then $x_m - x \in B_{1/n}(0)$. But that means that $\|x_m - x\| < 1/n$. Since $1/n \rightarrow 0$, then $\|x_n - x\| \rightarrow 0$ as well.

Example 1.2.2. Let's give some examples of Banach spaces.

- (1) \mathbb{R}^d and \mathbb{C}^d , as we learned in elementary real analysis.
- (2) $C([a, b])$ with $\|f\| = \sup_{x \in [a, b]} |f(x)|$ is Banach, because a sequence $\{f_n\}$ is Cauchy iff it converges uniformly, and we know the uniform limit of continuous functions is continuous.

$C([a, b])$ with norm

$$\|f\|_{L^1} = \int_a^b |f(x)| dx$$

is *not* complete, and therefore not Banach! This will verify the statement we made last lecture, that these spaces aren't the same. This is interesting behavior, because it doesn't happen in finite dimensions, and is an example of the subtle differences in behavior between finite-dimensional and infinite-dimensional vector spaces.

We'll let $a = -1$ and $b = 1$, though by suitable rescaling or translation this works for any $[a, b]$ with a and b finite.

Let $f_n(x)$ be 1 on $[-1, 0]$, then decrease linearly on $[0, 1/n]$, and then be 0 on $[1/n, 1]$. Then,

$$\begin{aligned} \|f_n - f_m\|_{L^1} &= \int_{-1}^1 |f_n(x) - f_m(x)| dx \\ &= \int_0^1 |f_n(x) - f_m(x)| dx \\ &\leq \int_0^1 (|f_n(x)| + |f_m(x)|) dx \\ &= \frac{1}{2n} + \frac{1}{2m}. \end{aligned}$$

This goes to 0, so $\{f_n\}$ is Cauchy. But it converges to the step function

$$f(x) = \begin{cases} 1, & x < 0 \\ 0, & x > 0. \end{cases}$$

This is because

$$\begin{aligned} \|f_n - f\|_{L^1} &= \int_{-1}^1 |f_n(x) - f(x)| dx \\ &= \int_0^1 |f_n(x)| dx = \frac{1}{2n}, \end{aligned}$$

which goes to 0, so $f_n \rightarrow f$ after all.

This means that when we talk about $C([a, b])$, unless otherwise specified, we'll use the other norm, which makes it into a Banach space.

This situation, where the same vector space has two norms with different topological properties, is actually fairly common.

Definition. Let X be a vector space and $\|\cdot\|_1$ and $\|\cdot\|_2$ be norms on X . One says that the two norms are *equivalent* if there exist $c, d > 0$ such that for all $x \in X$, $c\|x\|_1 \leq \|x\|_2 \leq d\|x\|_1$.

This means that, though they might not agree precisely, the vague notions of "small" and "large" are the same in both norms.

We'll see eventually that all norms on a finite-dimensional space are equivalent, even though we already know that $\|\cdot\|$ and $\|\cdot\|_{L^1}$ are inequivalent on $C([a, b])$. We do know, however, that for $f \in C([0, 1])$, $\|f\|_{L^1} \leq \|f\|$,³ but the other bound fails: there is no constant C such that $\|f\| \leq C\|f\|_{L^1}$. We'll see this using the sequence $\{f_n\}$, where f_n increases linearly from 0 to n on $[0, 1/n]$, decreases on $[1/n, 2/n]$, and is 0 elsewhere. This sweeps out a triangle, so $\|f_n\| = n$, but $\|f_n\|_{L^1} = 1$ for all n , and thus no such C exists.

Proposition 1.2.3. Let $\|\cdot\|_1$ and $\|\cdot\|_2$ be two equivalent norms on X . Then, their induced topologies are the same.

To be precise, the collections of open sets \mathcal{O}_1 and \mathcal{O}_2 induced from $\|\cdot\|_1$ and $\|\cdot\|_2$, respectively, are identical.

³More generally, on $C([a, b])$, $\|f\|_{L^1} \leq (b-a)\|f\|$.

PROOF. We'll let $B_r^1(x)$ denote the ball of radius r around x in $\|\cdot\|_1$, and define $B_r^2(x)$ similarly.

Since $\|\cdot\|_1$ and $\|\cdot\|_2$ are equivalent, there exist c and d such that for any x and r , $B_{r/d}^1(x) \subseteq B_r^2(x) \subseteq B_{r/c}^1(x)$. Thus, if O_2 is any open set in \mathcal{O}_2 , then for any $x \in O_2$, there's an r such that $B_r^2(x) \subseteq O_2$, and therefore $B_{r/d}^1(x) \subseteq O_2$, and so O_2 is open in \mathcal{O}_1 , and the argument in the other direction is similar. \square

Convexity. Convexity is an important notion because it allows us to talk about the line joining two points.

Definition. Let X be a vector space over \mathbb{F} . Then, a set $C \subseteq X$ is convex if whenever $x, y \in C$, the line $\{tx + (1-t)y : 0 \leq t \leq 1\}$ is contained in C .

Proposition 1.2.4. In any NLS, $B_r(x)$ is convex.

PROOF. Let $y, z \in B_r(x)$ and $t \in [0, 1]$. We want to show that $ty + (1-t)z \in B_r(x)$. We'll have to write x as $x + tx - tx$ and then use the triangle inequality. Specifically,

$$\begin{aligned} \|ty + (1-t)z - x\| &= \|t(y - x) + (1-t)(z - x)\| \\ &\leq t\|y - x\| + (1-t)\|z - x\| \\ &< tr + (1-t)r = r. \end{aligned} \quad \square$$

This is more interesting than it looks, because in some spaces that are otherwise similar to NLSes, there exist balls that are non-convex.

Even in finite dimensions, balls aren't necessarily round; they can even be square! But that doesn't make much of a difference.

Linear Operators. We'll talk about linear operators in order to manipulate and transform functions.

Definition. A linear operator is a function $T : X \rightarrow Y$ of vector spaces X and Y such that

- (1) $T(x + y) = T(x) + T(y)$, and
- (2) $T(\lambda x) = \lambda T(x)$.

The idea is that scalar multiplication and addition in X and Y (which are *a priori* very different) are considered the same by T , which commutes with them.

Definition. A linear operator $T : X \rightarrow Y$, where X and Y are NLSes, is *bounded* if it takes bounded sets to bounded sets.

That is, if $C \subseteq X$ is bounded, then $T(C) = \{y : y = T(x) \text{ for some } x \in C\}$.

The definition is nice, but everybody thinks of bounded operators by the following characterization.

Proposition 1.2.5. Let X and Y be normed linear spaces and $T : X \rightarrow Y$ be linear. Then, T is bounded iff there exists an $C > 0$ such that $\|Tx\|_Y \leq C\|x\|_X$ for all $x \in X$.

PROOF. First, suppose T is bounded. Then, the image of $B_1(0)$ (in X) is some bounded set, and therefore contained in a ball $B_R(0)$ for some R . In particular, if $y \in B_1(0)$, then $\|Ty\|_Y \leq R$.

Given $x \in X$, if $x = 0$ then $Tx = 0$, so we're good. If $x \neq 0$, let $y = (1/2\|x\|_X) \cdot x$, so that $\|y\| = 1/2$, and therefore $y \in B_1(0)$, and therefore $\|Ty\| \leq R$. That is,

$$\left\| T\left(\frac{1}{2\|x\|} \|x\|\right) \right\| = \frac{1}{2\|x\|} \|Tx\| \leq R,$$

and therefore $\|Tx\| \leq 2R\|x\|$, so with $C = 2R$ we're done.

Conversely, suppose there exists a $C > 0$ such that $\|Tx\| \leq C\|x\|$ for all $x \in X$. Let $M \subseteq X$ be bounded; then, $M \subseteq B_R(0)$ for some R . For an $x \in M$, $\|Tx\| \leq C\|x\| \leq CR$, so $T(X) \subseteq B_{CR}(0)$ in Y , and thus T is bounded. \square

Lecture 3.

Bounded Linear Operators: 8/31/15

Let X and Y be normed linear spaces; the maps between them that we'll consider are linear operators $T : X \rightarrow Y$, as in the previous lecture.

If T is one-to-one and onto, then we should have an inverse $T^{-1} : Y \rightarrow X$. It's easy to check that T^{-1} is linear; you probably checked this as an undergraduate. In this situation, we have structure preservation: it doesn't matter

whether you check addition in X or in Y , or scalar multiplication. Thus, in the sense of linear algebra, X and Y look the same; they have the same addition and scalar multiplication. In this case, we say that X and Y are *isomorphic*; they may be unequal as sets (e.g. sequences or functions), but identical from the perspective of linear algebra.

For vector spaces, these maps are pretty cool, but for topology, we care about continuous maps $f : X \rightarrow Y$. Thus, as you might guess, when studying normed linear spaces, we care about maps $X \rightarrow Y$ that are both linear and continuous.

Definition. If X and Y are NLSes, then $B(X, Y)$ denotes the set of functions $f : X \rightarrow Y$ that are both linear and continuous.

Continuity means that for all $\varepsilon > 0$ there exists a $\delta > 0$ depending on x and ε such that when $d(x, y) < \delta$, then $d(f(x), f(y)) \leq \varepsilon$. But since there's a norm defining the metric, this is equivalent to stating that when $\|x - y\| < \delta$, then $\|f(x) - f(y)\| \leq \varepsilon$. And if $f = T$ is a linear operator, then $\|T(x) - T(y)\| < \varepsilon$ is equivalent to requiring $\|T(x - y)\| \leq \varepsilon$. In other words, this doesn't depend on x at all: letting $z = x - y$, continuity of a linear $T : X \rightarrow Y$ means that when $\|z\| < \delta$, then $\|Tz\| \leq \varepsilon$.

In other words, if you know what a linear map does around 0, you know what it looks like everywhere.

Proposition 1.3.1. *Let X and Y be NLSes and $T : X \rightarrow Y$ be linear. Then, the following are equivalent:*

- (1) T is continuous.
- (2) T is continuous at some $x_0 \in X$.
- (3) T is bounded.

This is why we used the notation $B(X, Y)$: it stands for “bounded.” And we can now talk about bounded linear maps, with continuity understood.

PROOF. Clearly, (1) \implies (2). For (2) \implies (3), suppose T is continuous at some $x_0 \in X$. With $\varepsilon = 1$, this means there's a $\delta > 0$ such that $\|x - x_0\| \leq \delta$ implies $\|Tx - Tx_0\| \leq 1$, i.e. $\|T(x - x_0)\| \leq 1$. In other words, with $z = x - x_0$, when $\|z\| \leq \delta$, we have $\|Tz\| \leq 1$.

For $x = 0$ boundedness is clear, but if $x \neq 0$, then

$$\begin{aligned} \|Tx\|_Y &= \left\| \frac{\|x\|}{\delta} T\left(\frac{\delta x}{\|x\|}\right) \right\|_Y \\ &= \frac{\|x\|}{\delta} \left\| T\left(\frac{\delta x}{\|x\|}\right) \right\|_Y \leq \frac{1}{\delta} \|x\|_X, \end{aligned}$$

so with $C = 1/\delta$, T is a bounded operator.

For (3) \implies (1), we know $\|Tx\|_Y \leq C\|x\|_X$ for some fixed C and all $x \in X$. Let $\varepsilon > 0$ and pick any $x_0 \in X$. Then, if $\delta = \varepsilon/C$ and $\|x - x_0\| \leq \delta$, then

$$\|T(x - x_0)\| \leq C\|x - x_0\| \leq C\delta = \varepsilon,$$

so T is continuous at x_0 and therefore everywhere. \square

It turns out $B(X, Y)$ is a vector space itself, with $(f + g)(x) = f(x) + g(x)$ and $(\lambda \cdot f)(x) = \lambda \cdot (f(x))$, which is little surprise. But we do have to check that if $f = T$ and $g = S$ are linear, $f + g$ and λf are also linear, i.e. $(T + S)(x + y) = (T + S)(x) + (T + S)(y)$, and similarly for scalar multiplication.

What makes this more interesting is that $B(X, Y)$ is an NLS itself. What's the norm, you ask? Excellent question. The norm is

$$\|T\| = \|T\|_{B(X, Y)} = \sup_{x \in B_1(0)} \|Tx\|_Y.$$

Since T is continuous and bounded, $T(B_1(0))$ is a bounded set. Then, the norm of T is the radius of the smallest ball that contains $T(B_1(0))$, which is the supremum of the amount that T scales any point in the unit ball. Since T is bounded, the norm is a finite, nonnegative number.

Note that, even though we called this a norm, we still have to check that it's a norm!

Proposition 1.3.2. *Let X and Y be NLSes. Then, $\|\cdot\|_{B(X, Y)}$ is a norm on $B(X, Y)$. Moreover, if $T \in B(X, Y)$,*

$$\|T\| = \sup_{\|x\|_X \leq 1} \|Tx\|_Y = \sup_{\|x\|_X = 1} \|Tx\|_Y = \sup_{x \neq 0} \frac{\|Tx\|_X}{\|x\|_X}.$$

Furthermore, if Y is Banach, then $B(X, Y)$ is too.

This last point is quite interesting: completeness follows when the range is complete, but the domain doesn't matter.

PROOF. First, that $\|\cdot\|$ is a norm: we have three properties to show.

- We need $\|T\| = 0$ iff $T = 0$. Clearly, if $T = 0$ (i.e. $T(x) = 0$ for all x), then $\|T\| = \sup_{x \in B_1(0)} \|Tx\| = \|0\| = 0$. Conversely, if we assume $\|T\| = 0$, then for any $x \in B_1(0)$, $\|Tx\| = 0$, so $Tx = 0$. Thus, $T|_{B_1(0)} = 0$. For general x , we'll scale $x = 2\|x\|(x/2\|x\|)$, so

$$Tx = 2\|x\|T\left(\frac{x}{2\|x\|}\right) = 2\|x\| \cdot 0 = 0,$$

since $x/2\|x\| \in B_1(0)$. Thus, $T = 0$.

- For linearity of the norm,

$$\|\lambda T\| = \sup_{x \in B_1(0)} \|\lambda Tx\| = \sup_{x \in B_1(0)} |\lambda| \|Tx\| = |\lambda| \sup_{x \in B_1(0)} \|Tx\| = |\lambda| \|T\|.$$

Exercise. Finish the proof that this is a norm by addressing the triangle inequality, which isn't too complicated.

Next, we have the different ways of calculating the norm. The idea is that since T is continuous, the supremum shouldn't depend on whether the boundary is present or not. One interesting corollary of the formulas for calculating $\|T\|$ is that for any $x \in X$, $\|Tx\| \leq \|T\|\|x\|$.

The last part does require care. Let $\{T_n\}_{n=1}^\infty$ be a Cauchy sequence. That is, given an $\varepsilon > 0$, there's an $N > 0$ such that if $m, n \geq N$, then $\|T_n - T_m\|_{B(X,Y)} \leq \varepsilon$. Thus, given an $x \in X$, $\|T_n x - T_m x\|_Y \leq \|T_n - T_m\| \|x\|_X$. The right-hand side goes to 0 as a Cauchy sequence in m and n , and therefore the left-hand side does too. That is, $\{T_n x\}_{n=1}^\infty \subset Y$ is a Cauchy sequence. Since Y is Banach, this means there's a limit $\lim_{n \rightarrow \infty} T_n x = T(x) \in Y$. This defines a map $T : X \rightarrow Y$; we need to prove that it's bounded linear and that $T_n \rightarrow T$.

First, let's look at linearity.

$$T(x + y) = \lim_{n \rightarrow \infty} T_n(x + y) = \lim_{n \rightarrow \infty} (T_n x + T_n y).$$

Since addition is continuous, we can break this up as

$$= \lim_{n \rightarrow \infty} T_n x + \lim_{n \rightarrow \infty} T_n y = Tx + Ty.$$

Similarly, since scalar multiplication is continuous,

$$T(\lambda x) = \lim_{n \rightarrow \infty} T_n(\lambda x) = \lambda T(x).$$

Next, let's check that T is bounded. Since the norm is continuous,

$$\begin{aligned} \|Tx\|_Y &= \left\| \lim_{n \rightarrow \infty} T_n x \right\|_Y \\ &= \lim_{n \rightarrow \infty} \|T_n x\|_Y. \end{aligned}$$

However, this limit *a priori* might not exist, so we have to use the \limsup .

$$\begin{aligned} &\leq \limsup_{n \rightarrow \infty} \|T_n\| \|x\|_X \\ &= M \|x\|_X. \end{aligned}$$

Here, M is an upper bound on $\|T_n\|$, because $\{T_n\}$ is Cauchy and therefore bounded. Thus, we know $T \in B(X, Y)$.

Finally, to show $T_n \rightarrow T$, we need to be careful: limits depend on the topology that we're using, and so we should be careful that we're using the topology defined by $\|\cdot\|_{B(X,Y)}$.

Let $x \in B_1(0)$. Then,

$$\begin{aligned} \|Tx - Ty\|_Y &= \lim_{m \rightarrow \infty} \|T_m x - T_n x\| \\ &= \lim_{m \rightarrow \infty} \|(T_m - T_n)x\| \\ &\leq \limsup_{m \rightarrow \infty} \|T_m - T_n\| \|x\|. \end{aligned}$$

Since $\{T_n\}$ is Cauchy, then for any $\varepsilon > 0$, $\|T_m - T_n\| \leq \varepsilon$ when m, n are sufficiently large, and therefore the \limsup goes to 0 as $n \rightarrow \infty$, and so $T_n \rightarrow T$. \square

There's one particularly important case, in which $Y = \mathbb{F}$.

Definition. The *dual space* of an NLS X is $X^* = B(X, \mathbb{F})$.

By Proposition 1.3.2, X^* is always a Banach space.

Though $B(X, Y)$ can be complicated for general Y , one can often understand it more easily using X^* .

Example 1.3.3. We can connect this with finite-dimensional linear algebra that we're more familiar with, and see that it's actually quite special.

Let X be a d -dimensional vector space over \mathbb{F} with basis $\{e_n\}_{n=1}^d$. Thus,

$$\begin{aligned} X &= \text{span}\{e_1, \dots, e_d\} \\ &= \{\alpha_1 e_1 + \dots + \alpha_d e_d \mid \alpha_i \in \mathbb{F}\}, \end{aligned}$$

and we can write $x = x_1 e_1 + \dots + x_d e_d \in X$. The map $T : X \rightarrow \mathbb{F}^d$ sending $x \mapsto (x_1, \dots, x_d)$ is one-to-one, onto, and linear, so all finite-dimensional vector spaces over a specified field are isomorphic. Moreover, we showed that all norms over a finite-dimensional vector space are equivalent, so as NLSes, they're all isomorphic too! There are many norms, which may still be interesting, but there's only one topology.

Lecture 4.

ℓ^p -norms: 9/2/15

Recall that we were looking at examples of Banach spaces, and that the first examples we saw (Example 1.3.3) were finite-dimensional vector spaces. If $d = \dim X$ is finite, so that $X = \text{span}\{e_1, \dots, e_n\}$ (which is a basis for X), then the map $T : X \rightarrow \mathbb{F}^d$ sending $(x_1 e_1 + \dots + x_d e_d) \mapsto (x_1, \dots, x_d)$ is an isomorphism of vector spaces, and the claim is that these maps define the same topology as well.

But first, let's define some norms on \mathbb{F}^d . Let $1 \leq p \leq \infty$, and define

$$\|x\|_{\ell^p} = \begin{cases} \left(\sum_{n=1}^d |x_n|^p \right)^{1/p}, & p < \infty \\ \max_n |x_n|, & p = \infty. \end{cases}$$

Sometimes, these are denoted $\|x\|_{\ell_p}$. Also, the case $p = 2$ is our familiar Euclidean norm $\|x\|_{\ell_2} = |x|$.

We do have to show that these are norms. When $p = 1, \infty$, it's a straightforward check, and when $1 < p < \infty$, the first two properties are pretty simple, but the triangle inequality is harder.

Lemma 1.4.1 (Young's inequality⁴). *Let $1 < p < \infty$ and q be the conjugate exponent defined such that $1/p + 1/q = 1$. If $a, b \geq 0$, then $ab \leq a^p/p + b^q/q$, with equality iff $a^p = b^q$. Moreover, for all $\varepsilon > 0$, there exists a C depending on p and ε such that $ab \leq \varepsilon a^p + C b^q$.*

PROOF. The proof is easy once you know the trick, to look at the right function. Let $u : [0, \infty) \rightarrow \mathbb{R}$ send

$$u(t) = \frac{t^p}{p} + \frac{1}{q} - t.$$

Its derivative is well-defined: $u'(t) = t^{p-1} - 1$, so $u'(0) = 1$. In particular, $u(0) = 1/q$, and $u(1) = 0$ is a strict minimum.

We'll apply this to $t = ab^{-q/p}$:

$$\begin{aligned} 0 \leq u(ab^{-q/p}) &= \frac{a^p}{p b^q} + \frac{1}{q} - \frac{a}{b^{q/p}} \\ &= \frac{1}{b^q} \left(\frac{a^p}{p} + \frac{b^q}{q} - \frac{a b^q}{b^{q/p}} \right), \end{aligned}$$

but $b^q/b^{q/p} = b$, since $q - q/p = q(1 - 1/p) = 1$. Thus, $0 \leq a^p/p + b^q/q - ab$, and equality holds iff $t = ab^{-q/p} = 1$, where $u(t)$ is equal to 0.

For the second part, we can write

$$ab = ((\varepsilon p)^{1/p} a)((\varepsilon p)^{-1/p} b) \leq \frac{\varepsilon p a^p}{p} + \frac{(\varepsilon p)^{-q/p}}{q} b^q. \quad \square$$

⁴Young's inequality technically refers to a more general statement; this could be called "Young's inequality for products."

For conjugate exponents, we have the convention that the conjugate of 1 is ∞ , and vice versa.

Theorem 1.4.2 (Hölder's inequality). *Let $1 \leq p \leq \infty$ and q be its conjugate exponent. If $x, y \in \mathbb{F}^d$, then*

$$\sum_n |x_n y_n| \leq \|x\|_{\ell^p} \|y\|_{\ell^q}.$$

When $p = 2$, this is also known as the *Cauchy-Schwarz inequality*.

PROOF. The cases $p = 1, \infty$ are trivial; expand their definitions out. Similarly, if $x = 0$ or $y = 0$, there's not a lot to say. Thus, we're left with $1 < p < \infty$, so we can use Lemma 1.4.1.

Let $a = |x_n|/\|x\|_{\ell^p}$ and $b = |y_n|/\|y\|_{\ell^q}$. Then, by Lemma 1.4.1,

$$\frac{|x_n|}{\|x\|_{\ell^p}} \frac{|y_n|}{\|y\|_{\ell^q}} \leq \frac{|x_n|^p}{p\|x\|_{\ell^p}^p} + \frac{|y_n|^q}{q\|y\|_{\ell^q}^q},$$

so summing all n of those,

$$\begin{aligned} \frac{\sum_n |x_n y_n|}{\|x\|_{\ell^p} \|y\|_{\ell^q}} &\leq \frac{\sum_n |x_n|^p}{p\|x\|_{\ell^p}^p} + \frac{\sum_n |y_n|^q}{q\|y\|_{\ell^q}^q} \\ &= \frac{\|x\|_{\ell^p}^p}{p\|x\|_{\ell^p}^p} + \frac{\|y\|_{\ell^q}^q}{q\|y\|_{\ell^q}^q} \\ &= \frac{1}{p} + \frac{1}{q} = 1. \end{aligned} \quad \square$$

Now, we can use this to prove the triangle inequality for $\|\cdot\|_{\ell^p}$. We'll need two things for the Hölder inequality, so just take one term out of the p^{th} power:

$$\begin{aligned} \|x + y\|_{\ell^p}^p &= \sum_{n=1}^d |x_n + y_n|^p \\ &\leq \sum_{n=1}^d |x_n + y_n|^{p-1} (|x_n| + |y_n|) \\ &\leq \left(\sum_{n=1}^d |x_n + y_n|^{(p-1)q} \right)^{1/q} (\|x\|_{\ell^p} + \|y\|_{\ell^q}). \end{aligned}$$

Since p and q are conjugate, $p = (p-1)q$, so the first term is $\|x + y\|_{\ell^p}^{p/q}$. Thus,

$$\|x + y\|_{\ell^p}^{p-p/q} \leq \|x\|_{\ell^p} + \|y\|_{\ell^p},$$

and $p - p/q = 1$, so we're done.

Moreover, all these norms are equivalent.

Proposition 1.4.3. *Let $1 \leq p \leq \infty$. Then, for all $x \in \mathbb{F}^d$,*

$$\|x\|_{\ell^\infty} \leq \|x\|_{\ell^p} \leq d^{1/p} \|x\|_{\ell^\infty}.$$

These estimates are sharp, the first at $x = (1, 0, 0, \dots, 0)$, and the second at $x = (1, 1, \dots, 1)$.

PROOF. Let m be an index for which $|x_m| = \max_n |x_n|$. Since $f(x) = x^{1/p}$ is an increasing function,

$$\|x\|_{\ell^\infty} = |x_m| = (|x_m|^p)^{1/p} \leq \left(\sum_{n=1}^d |x_n|^p \right)^{1/p} = \|x\|_{\ell^p},$$

and

$$\begin{aligned}\|x\|_{\ell^p} &= \left(\sum_{n=1}^d |x_n|^p \right)^{1/p} \\ &\leq \left(\sum_{m=1}^d |x_m|^p \right)^{1/p} \\ &= (d|x_m|^p)^{1/p} = d^{1/p} \|x\|_{\ell^\infty}.\end{aligned}\quad \square$$

Notice that some of these proof methods fail horribly in infinite dimensions.

It turns out that on all finite-dimensional vector spaces, all norms are equivalent.

Proposition 1.4.4. *All norms on a finite-dimensional NLS are equivalent. Moreover, a $K \subset X$ is compact iff it is closed and bounded.*

That means there's only one topology.

PROOF. Let $d = \dim X$ and $\{e_n\}_{n=1}^d$ be a basis. Then, let $T : X \rightarrow \mathbb{F}^d$ be the coordinate map defined above. Let \cong denote an isomorphism of NLSes.

We'll define a norm $\|\cdot\|_1$ on x by $\|x\|_1 = \|Tx\|_{\ell^1}$: of the three properties, the last two are trivial (since T is linear), so we just need to prove that $\|x\|_1 = 0$ iff $x = 0$. But T is one-to-one and onto, so this follows, and $\|\cdot\|_1$ is in fact a norm.⁵

Thus, $(X, \|\cdot\|_1) \cong (\mathbb{F}^d, \|\cdot\|_{\ell^1})$, so they really are the “same” space. This is because $T : X \rightarrow \mathbb{F}^d$ is a bounded map, with $C = 1$, and therefore continuous, and T^{-1} is also linear and continuous. Thus, T is an isomorphism of vector spaces and a homeomorphism of topological spaces, so we can take results in \mathbb{F}^d and apply them to X .

The Heine-Borel theorem from undergraduate real analysis tells us that $K \subset \mathbb{F}^d$ is closed and bounded iff it's compact. But since X and \mathbb{F}^d have the same topology, then this is also true in X . In particular, $S_1^1 = \{x \in X : \|x\|_1 = 1\}$ is also compact.

Now, for any norm $\|\cdot\|$ on X and $x \in X$,

$$\|x\| = \left\| \sum_{n=1}^d x_n e_n \right\| \leq \sum_{n=1}^d |x_n| \|e_n\| \leq C \|x\|_1,$$

where $C = \max_n \|e_n\|$. Notice that this step won't work in infinite dimensions. Our upper bound implies that $(\text{Top})_{\|\cdot\|} \subseteq (\text{Top})_{\|\cdot\|_1}$, so the former topology is said to be stronger. We'll prove the two are equal by providing a lower bound.

We have a continuous map $\|\cdot\| : (X, \|\cdot\|_1) \rightarrow \mathbb{R}$. It's also continuous as a map $\|\cdot\| : (X, \|\cdot\|) \rightarrow \mathbb{R}$. Let $a = \inf_{x \in S_1^1} \|x\|$; since S_1^1 is compact and the norm is continuous, the minimum is attained, and it must be positive (because $0 \notin S_1^1$).

Thus, for any $x \in X$, $\|x\|/\|x\|_1 \geq a$, so $\|x\| \geq a\|x\|_1$, which is our desired lower bound. □

Corollary 1.4.5. *If X is a d -dimensional NLS, then $X \cong \mathbb{F}^d$.*

Corollary 1.4.6. *If X and Y are NLSes and X is finite-dimensional, then every linear $T : X \rightarrow Y$ is bounded and $X^* = \mathbb{F}^d$, given by $T(x) = y \cdot x$.*

Lecture 5.

ℓ^p and L^p -spaces: 9/4/15

“There are different sizes of infinity, and this one is the best.”

Last time we showed that if $(X, \|\cdot\|)$ is a finite-dimensional NLS, then it's isomorphic and homeomorphic to $(\mathbb{F}^d, \|\cdot\|_{\ell^2})$, where $d = \dim X$. Moreover, X is Banach, and $(\mathbb{F}^d)^* \cong \mathbb{F}^d$. Finite dimensions aren't very interesting, but they're a good place to gain intuition.

A lot of this nice stuff goes away for infinite-dimensional spaces, and some are nicer than others.

⁵A great way to create a new norm is to map from one space to another (or the same one) and pull the norm back.

Example 1.5.1. Let $1 \leq p \leq \infty$. We'll define a space ℓ^p which behaves sort of like an " \mathbb{F}^∞ ." Specifically,

$$\ell^p = \{x = \{x_n\}_{n=1}^\infty : x_n \in \mathbb{F}, \|x\|_{\ell^p} < \infty\},$$

where

$$\|x\|_{\ell^p} = \begin{cases} \left(\sum_{n=1}^\infty |x_n|^p \right)^{1/p}, & p \text{ finite} \\ \sup_n |x_n|, & p = \infty. \end{cases}$$

The same proofs for the ℓ^p -norms in finite-dimensional spaces apply, and show that ℓ^p is an NLS.

Theorem 1.5.2 (Hölder's inequality in ℓ^p). If $1 \leq p \leq \infty$, $1/p + 1/q = 1$, and $x \in \ell^p$ and $y \in \ell^q$, then

$$\sum_{n=1}^\infty |x_n y_n| \leq \|x\|_{\ell^p} \|y\|_{\ell^q}.$$

Again, the proof is identical to the one for the finite-dimensional ℓ^p -norm.

Note that ℓ^∞ can be a bit weird relative to the rest of the ℓ^p spaces.

If p is finite, then ℓ^p has countably infinite dimension, i.e. it has a basis that's countable. This is subtle: the span of a basis is the set of *finite* linear combinations; in the infinite case, we would have to worry about convergence. Anyways, set

$$e^{i_n} = \begin{cases} 1, & i = n \\ 0, & i \neq n. \end{cases}$$

Then, a basis for ℓ^p , called the *Schauder basis*, is $\mathcal{B} = \{e^i\}_{i=1}^\infty$, and its span is

$$\text{span}(\mathcal{B}) = \{\alpha_{i_1} e^{i_1} + \alpha_{i_2} e^{i_2} + \cdots + \alpha_{i_n} e^{i_n} : n \in \mathbb{N}, \alpha_{i_j} \in \mathbb{F}\}.$$

Note that this is *not* a basis in the linear-algebraic sense (which would have to be uncountable); rather, this means that ℓ^p is the closure of $\text{span}(\mathcal{B})$. That is, for all $x \in \ell^p$, there's a unique representation $x = \sum_{j=1}^\infty x_j e^j$, meaning that if x_N denotes the N^{th} partial sum, then $x_N \in \mathcal{B}$ for all N , and

$$\|x - x_N\|_{\ell^p} = \left(\sum_{n=N+1}^\infty |x_n|^p \right)^{1/p} \rightarrow 0.$$

This is a little weird, but the point is that, since you can't take infinite sums in a basis, things can get a little strange. But everything comes from the finite case.

ℓ^∞ does *not* have a countable basis. As a result, we sometimes consider subspaces with a countable basis. Define

$$c_0 = \{x \in \ell^\infty : \lim_{n \rightarrow \infty} x_n = 0\} \text{ and}$$

$$f_0 = \{x \in \ell^\infty : x_n = 0 \text{ for all but finitely many } n\}.$$

For example, $(1, 1, 1, \dots) \in \ell^\infty$, but it's not in c_0 or f_0 , and $(1, 1/2, 1/3, \dots)$ is in c_0 but not f_0 . f_0 and c_0 inherit the ℓ^∞ -norm and become NLSes in their own right.

If $1 \leq p \leq q < \infty$, then we have the following chain of inclusions:

$$f_0 \subseteq \ell^p \subseteq \ell^q \subseteq c_0 \subseteq \ell^\infty.$$

If you're looking for examples (or, sometimes, counterexamples), c_0 and f_0 are often useful. For example, on f_0 , we have a function $T : f_0 \rightarrow \mathbb{F}$ defined by

$$T(x) = \sum_{n=1}^\infty n x_n.$$

Since each $\alpha \in f_0$ is a finite sequence, then this is well-defined, and it's linear, but it's not bounded, since $T(e^i) = i$ but $\|e^i\|_{\ell^\infty} = 1$ for all i . Thus, we have a linear map which is not continuous.

Exercise. If $1 \leq p \leq \infty$, show that ℓ^p is Banach.

This is conceptually easy but a bit of work, coming down to calculus, and so we know that limits of Cauchy sequences exist. However, since ℓ^1 is a subspace of ℓ^∞ , we can consider the NLS $(\ell^1, \|\cdot\|_{\ell^\infty})$; this space is not Banach.

Lemma 1.5.3. Let $0 < p < 1$ and define ℓ^p in the same way as above. In this case, however, ℓ^p is not an NLS, because $\|\cdot\|_{\ell^p}$ isn't a norm.

PROOF. We can look at $(\mathbb{F}^2, \|\cdot\|_{\ell^p})$ to see this: we proved that, given the triangle inequality, the unit ball is convex. However, the unit ball isn't convex when $p < 1$. \square

The Hölder inequality allows us to create many continuous linear functionals $T : \ell^p \rightarrow \mathbb{F}$ when $1 \leq p \leq \infty$. Let q be the conjugate exponent (so $1/p + 1/q = 1$), and choose any $y \in \ell^q$. Then, we can produce a $T_y \in (\ell^p)^*$, i.e. $T_y : \ell^p \rightarrow \mathbb{F}$, defined by

$$T_y(x) = \sum_{n=1}^{\infty} x_n y_n.$$

Moreover, T_y is bounded, because $|T_y(x)| \leq \|y\|_{\ell^q} \|x\|_{\ell^p}$.

This defines an inclusion $\ell^q \hookrightarrow (\ell^p)^*$.

Exercise. In fact, when p is finite, $\ell^q = (\ell^p)^*$. Moreover, $T : \ell^q \rightarrow (\ell^p)^*$ sending $T(y) \rightarrow T_y$ is a bounded operator, as $\|T_y\|_{(\ell^p)^*} = \|y\|_{\ell^q}$.

That is, the dual space is the conjugate space; to show this, figure out how to write $T(e^i)$ as y_i for some $y_i \in \ell^q$.

The above result is untrue for ℓ^∞ ; in fact, $(\ell^\infty)^* \supsetneq \ell^1$, but $c_0^* = \ell^1$.

That's all that we really need to say about ℓ^p for now; it's one step up from finite-dimensional spaces, and is a bit different, but not all that exotic. Right now, our examples are \mathbb{F}^d , which is finite-dimensional; ℓ^p when p is finite, which has countable dimension, and ℓ^∞ , which has uncountable dimension.

Lebesgue spaces. Let $\Omega \subseteq \mathbb{R}^d$ be a measurable set. We want to define a space of functions on Ω . However, when we talk about functions and measure, we really want to define two functions f and g as “the same” if $f(x) = g(x)$ except on a set of measure zero. If this is true, no integral can distinguish f and g .

Definition. Let $1 \leq p < \infty$, and define $L^p(\Omega)$ be the set of measurable functions⁶ $f : \Omega \rightarrow \mathbb{F}$ such that $\int_{\Omega} |f(x)|^p dx$ is finite. $L^p(\Omega)$ becomes an NLS with the norm

$$\|f\|_p = \left(\int_{\Omega} |f(x)|^p \right)^{1/p},$$

though we'll have to show that.

Once again, we can define this for $p < 1$, but it won't end up being a norm.

When $p = \infty$, we'll do things a little differently, as usual.

Definition.

- A measurable $f : \Omega \rightarrow \mathbb{F}$ is *essentially bounded* by $K \in \mathbb{R}$ if $|f(x)| \leq K$ for almost every $x \in \Omega$ (i.e. the set where this is not true has measure zero).
- The *essential supremum* of f , denoted $\text{ess sup}_{x \in \Omega} |f(x)|$, is the infimum of the K that essentially bound f .

Then, we can define $L^\infty(\Omega)$ as the set of (equivalence classes of) measurable functions whose essential suprema are finite, and $\|f\|_\infty = \text{ess sup}_{x \in \Omega} |f(x)|$. This will also be an NLS, though we'll have to show that too.

Proposition 1.5.4. If $0 < p \leq \infty$, then $L^p(\Omega)$ is a vector space, and $\|f\|_p = 0$ iff $f = 0$ almost everywhere on Ω .

PROOF. First, why is $L^p(\Omega)$ closed under addition? If p is finite, then

$$|f(x) + g(x)|^p \leq (|f(x)| + |g(x)|)^p \leq 2^p(|f(x)|^p + |g(x)|^p),$$

so when one integrates, if $f, g \in L^p(\Omega)$, then the rightmost quantity is bounded and therefore the leftmost one is. Scalar multiplication (and the scaling property of the norm) is easy: just write down the definition.

For $p = \infty$, the maximum of the sum cannot be bigger than the sum of the maxima, so $\|f + g\|_\infty = \|f\|_\infty + \|g\|_\infty$. Scaling and scalar multiplication are also straightforward. \square

Thus, all we have left is the triangle inequality, which we'll show next class.

⁶To be pedantic, the elements of $L^p(\Omega)$ are equivalence classes of functions that differ from f on a set of measure zero, since the integrals are the same.

Lecture 6.

 $L^p(\Omega)$ is Banach: 9/9/15

Recall that if $\Omega \subseteq \mathbb{R}^d$, then $L^p(\Omega)$ is the set of equivalence classes of measurable functions $\Omega \rightarrow \mathbb{F}$ with $\|f\|_p < \infty$, where $f \sim g$ if they differ on a set of measure zero. Then, the p -norm is

$$\|f\|_p = \begin{cases} \left(\int_{\Omega} |f(x)|^p dx \right)^{1/p}, & p < \infty \\ \text{ess sup}_{x \in \Omega} |f(x)|, & p = \infty. \end{cases}$$

Last time, we showed that $L^p(\Omega)$ is a vector space, and two of the properties of NLSes, the zero and scaling properties. Today we'll attack the triangle inequality; just as for ℓ^p , we'll need Hölder's inequality.

Proposition 1.6.1 (Hölder's inequality for L^p). *Let $1 \leq p \leq \infty$ and $1/p + 1/q = 1$. If $f \in L^p(\Omega)$ and $g \in L^q(\Omega)$, then $fg \in L^1(\Omega)$ and $\|fg\|_1 \leq \|f\|_p \|g\|_q$, with equality iff $|f(x)|^p$ is proportional to $|g(x)|^q$.*

PROOF. if $p = 1, \infty$, we already know that $\int_{\Omega} |f(x)g(x)| dx \leq \|g\|_{\infty} \int_{\Omega} |f| dx = \|f\|_1 \|g\|_{\infty}$.

If $1 < p < \infty$, we know from Lemma 1.4.1 that $ab \leq a^p/p + b^q/q$, with equality when $a^p = b^q$. If $\|f\|_p = 0$ or $\|g\|_q = 0$, then we're done; otherwise,

$$\frac{|f(x)|}{\|f\|_p} \frac{|g(x)|}{\|g\|_q} \leq \frac{|f(x)|^p}{\|f\|_p^p} + \frac{|g(x)|^q}{\|g\|_q^q},$$

so integrating, we get

$$\frac{\int |fg|}{\|f\|_p \|g\|_q} \leq 1,$$

with equality when $|f(x)|^p / \|f\|_p^p = |g(x)|^q / \|g\|_q^q$, which gives us our proportionality. ⊠

Theorem 1.6.2 (Minkowski's inequality). *If $1 \leq p \leq \infty$, then $\|f + g\|_p \leq \|f\|_p + \|g\|_p$.*

PROOF. Notice that if f or g isn't in $L^p(\Omega)$, then its p -norm is infinite, so we're done. The result is also clear if $p = 1, \infty$: the supremum of the sum is less than the sum of the suprema, and similarly with absolute value.

So we only have to worry about $1 < p < \infty$, and here we'll use a similar trick as for ℓ^p spaces, taking one copy of a p^{th} power.

$$\begin{aligned} \|f + g\|_p^p &= \int_{\Omega} |f(x) + g(x)|^p dx \\ &\leq \int_{\Omega} |f(x) + g(x)|^{p-1} (|f(x)| + |g(x)|) dx. \end{aligned}$$

Using Hölder's inequality,

$$\begin{aligned} &\leq \left(\int_{\Omega} |f(x) + g(x)|^{(p-1)q} \right)^{1/q} (\|f\|_p + \|g\|_p) \\ &= \|f + g\|_p^{p-1} (\|f\|_p + \|g\|_p), \end{aligned}$$

so dividing by $\|f + g\|_p^{p-1}$, we're done. ⊠

L^p spaces are very important in analysis, and form an important set of examples for NLSes. A little later, we'll show that they're complete, but we should note that we're measuring the size of a function using varying p , which measure different things, between emphasizing large values at a point, or large values at infinity.

On \mathbb{R} , imagine a function that goes to ∞ as $x \rightarrow 0^+$ and 0 as $x \rightarrow \infty$. If p is large, we're emphasizing the large values of the function, so if it grows too quickly it might not be in $L^p(\mathbb{R})$. If p is small, then we're emphasizing the long tail as $x \rightarrow \infty$; if it dies too slowly, it might not be in $L^p(\mathbb{R})$. An instructive example is x^p , which is in some L^q spaces but not others.

An easier way to think about this is to bound Ω , so we don't have to worry about long tails.

Proposition 1.6.3. *Let μ denote the Lebesgue measure, and suppose $\mu(\Omega)$ is finite. Let $1 \leq p \leq q \leq \infty$.*

(1) *If $f \in L^q(\Omega)$, then $f \in L^p(\Omega)$, and in fact $\|f\|_p = (\mu(\Omega))^{1/p-1/q} \|f\|_q$.*

- (2) If $f \in L^\infty(\Omega)$, then $f \in L^p(\Omega)$ for $1 \leq p \leq \infty$, and $\lim_{p \rightarrow \infty} \|f\|_p = \|f\|_\infty$.
 (3) If $f \in L^p(\Omega)$ for $1 \leq p < \infty$ and $\|f\|_p \leq K$ for all such p , then $f \in L^\infty(\Omega)$ and $\|f\|_\infty \leq K$.

These will be proven in the homework. Part (2) is the reason the L^∞ -norm is named such. Note also that there exist f such that $f \in L^p(\Omega)$ for $1 \leq p < \infty$ but $f \notin L^\infty(\Omega)$, even when Ω has finite measure.

The general proof idea is to consider sets of bad points and see what happens.

Proposition 1.6.4. For $1 \leq p \leq \infty$ and Ω measurable, $L^p(\Omega)$ is complete.

Thus, we have another useful class of Banach spaces.

PROOF. As usual, we'll start with a Cauchy sequence $\{f_n(x)\}_{n=1}^\infty$ in $L^p(\Omega)$. The idea will be to write

$$f_n(x) = f_1(x) + f_2(x) - f_1(x) + f_3(x) - f_2(x) + \cdots + f_n(x) - f_{n-1}(x),$$

so if we group the $f_i(x) - f_{i-1}(x)$, then these pieces should be small, and therefore we ought to converge to some function $f(x)$. There are technical problems, though, since we don't know how fast the f_n converge, so we need to try $f_i(x) - f_{i-k}(x)$ for $k > 1$. Moreover, we'll use absolute values. This is the idea; now, let's write it down carefully.

First, select a subsequence such that $\|f_{n_{j+1}} - f_{n_j}\| \leq 2^{-j}$ for all j ; we can do this because if we have n_{j-1} , there's an n_j such that $\|f_{n_j} - f_m\| \leq 2^{-j}$ when $m \geq n_j \geq n_{j-1}$.

Let

$$F_m(x) = |f_{n_1}(x)| + \sum_{j=1}^m |f_{n_{j+1}}(x) - f_{n_j}(x)| \geq 0,$$

and additionally $\{F_m(x)\}$ is an increasing function, so there's a limit (which might be ∞ , but that's OK). Let $F(x) = \lim_{m \rightarrow \infty} F_m(x) \in [0, \infty]$. Then,

$$\|F_m\|_p \leq \|f_{n_1}\|_p + \sum_{j=1}^m 2^{-j} \leq \|f_{n_1}\|_p + 1,$$

which is finite. But more interestingly, $F \in L^p(\Omega)$ too! We'll have to treat L^∞ as a special case again.

If p is finite, we'll use the monotone convergence theorem.

$$\begin{aligned} \int_{\Omega} |F(x)|^p dx &= \int_{\Omega} \lim_{m \rightarrow \infty} |F_m(x)|^p dx \\ &\leq \lim_{m \rightarrow \infty} \int_{\Omega} |F_m(x)|^p dx \\ &\leq \|f_{n_1}\|_p^p + 1, \end{aligned}$$

which is finite.

When $p = \infty$, then $|F_m(x)| \leq \|F_m\|_\infty \leq \|f_{n_1}\|_\infty + 1$ for all $x \notin A_m$, where $\mu(A_m) = 0$. Thus, if $A = \bigcup_{n=1}^\infty A_n$, then $\mu(A) = 0$ too. Thus, $|F(x)| = \lim_{m \rightarrow \infty} |F_m(x)| \leq K$ for some K and all $m, x \notin A$, so $F \in L^\infty(\Omega)$.

Now,

$$f_{n_{j+1}}(x) = f_{n_1}(x) + (f_{n_2}(x) - f_{n_1}(x)) + \cdots + (f_{n_{j+1}}(x) - f_{n_j}(x)).$$

Thus, this converges absolutely *pointwise*⁷ to some $f(x)$, so f is measurable. Now, $|f_{n_j}(x)| \leq F(x)$, so $|f(x)| \leq F(x)$, and therefore $f \in L^p(\Omega)$.

But we need that $\|f_{n_j} - f\|_p \rightarrow 0$, so let's think about that. Again, we have to argue differently when $p = \infty$.

When p is finite, we'll use the dominated convergence theorem on $|f_{n_j}(x) - f(x)| \leq F(x) + |f(x)| \in L^p(\Omega)$:

$$\lim_{j \rightarrow \infty} \int_{\Omega} |f_{n_j}(x) - f(x)|^p dx \leq \int_{\Omega} \lim_{j \rightarrow \infty} |f_{n_j}(x) - f(x)|^p dx \rightarrow 0.$$

When p is infinite, for any j and k , there's a set B_{n_j, n_k} with measure zero such that on $\Omega \setminus B_{n_j, n_k}$, $|f_{n_j}(x) - f_{n_k}(x)| \leq \|f_{n_j} - f_{n_k}\|_\infty$. Thus,

$$B = \bigcup_j \bigcup_k B_{n_j, n_k}$$

⁷We have multiple notions of convergence floating around; be careful to distinguish pointwise convergence, uniform convergence, and convergence in L^p .

is a countable union, so $\mu(B) = 0$. Since $\{f_{n_j}\}$ is Cauchy, then for any $x \notin B$ and $\varepsilon > 0$, there's an $N > 0$ such that if $j, k \geq N$, then $|f_{n_j}(x) - f_{n_k}(x)| \leq \varepsilon$, so taking the pointwise limit $f_k(x) \rightarrow f(x)$, $|f_{n_j}(x) - f(x)| \leq \varepsilon$. Thus, since we're avoiding B , $\|f_{n_j} - f\|_\infty \leq \varepsilon$.

We're almost done: we have $f_{n_j} \rightarrow f$ in L^p , but we need $f_n \rightarrow f$ in L^p . If $\varepsilon > 0$, then there exists an $N > 0$ such that $\|f_n - f_{n_j}\|_p \leq \varepsilon/2$ for all $n, n_j \geq N$. Therefore $|f_{n_j} - f|_p \leq \varepsilon/2$ for all $n_j \geq N$, and therefore the triangle inequality tells us that

$$\|f_n - f\|_p \leq \|f_n - f_{n_j}\|_p + \|f_{n_j} - f\|_p \leq \varepsilon. \quad \square$$

If you examine the proof, we've also proven an interesting result.

Corollary 1.6.5. *If $1 \leq p < \infty$ and $\{f_n\}_{n=1}^\infty$ is a sequence in $L^p(\Omega)$ converging to f in the L^p -norm, then there exists a subsequence $\{f_{n_j}\}_{j=1}^\infty$ such that $f_{n_j}(x) \rightarrow f(x)$ pointwise a.e.*

So convergence in L^p implies pointwise convergence of a subsequence almost everywhere. We'll use this later.

It turns out that the dual space to $L^p(\Omega)$ is $L^q(\Omega)$, where q is the conjugate exponent. Given a $g \in L^q(\Omega)$, define an operator $T_g : L^p(\Omega) \rightarrow \mathbb{R}$ by

$$T_g(f) = \int_{\Omega} f(x)g(x) dx,$$

which makes sense and is finite by Proposition 1.6.1. Thus, this is well-defined, and linear because the integral is. It's continuous, because it's bounded (by Hölder's inequality again): $T_g(f) \leq \|g\|_q \|f\|_p$, so $\|t_g\| \leq \|g\|_q$, and it's probably not a surprise that's actually an equality: choose something like $f(x) = |g(x)|^{q/p} / \|g\|_q$ (maybe with a power in the denominator), to see that the bound is sharp.

Thus, we've shown that $L^q(\Omega) \subseteq (L^p(\Omega))^*$ in some sense, for $1 \leq p \leq \infty$. However, if p is finite, then $L^q(\Omega) = (L^p(\Omega))^*$; there are no other continuous linear functionals. When $p = \infty$, there are more, so the dual space is the space of positive measures: $g(x) dx$ is a measure, but there are other measures that aren't of that form.

We won't prove this, but it follows from a deep theorem in analysis called the Radon-Nikodym theorem.

Lecture 7.

The Hahn-Banach Theorem: 9/11/15

"Almost everything has three properties. Have you noticed that?"

Corollary 1.7.1. *Let X be an NLS, $Y \subset X$ be a linear subspace, and $f : Y \rightarrow \mathbb{R}$ be bounded. Then, there exists an $F \in X^*$ such that $F|_Y = f$ and $\|F\|_{X^*} = \|f\|_{Y^*}$.*

Though L^p functions can be complicated, all of them can be well-approximated by less complicated functions. Recall that a *simple* function is a Lebesgue-integrable function that takes on only finitely many values, and that a function is *compactly supported* if it is equal to 0 outside of a compact set.

Proposition 1.7.2. *For $1 \leq p \leq \infty$, the set \mathcal{S} of all measurable simple functions with compact support is dense in $L^p(\Omega)$.*

This says that for any $f \in L^p(\Omega)$ and $\varepsilon > 0$, there's a $\varphi \in \mathcal{S}$ such that $\|f - \varphi\|_{L^p(\Omega)} < \varepsilon$. The proof comes from measure theory: the integral was defined by the limit of approximations by simple functions, and so the integrals are successively better approximations.

Definition. Let $C_0^\infty(\Omega)$ denote the space of compactly supported, continuous functions.

Proposition 1.7.3. *If Ω is an open set and $1 \leq p < \infty$, then $C_0^\infty(\Omega)$ is dense in $L^p(\Omega)$.*

The proof follows from another measure-theoretic result called Lusin's theorem.

Now, we'll move into some deeper (and, well, harder) theorems and questions in functional analysis. We'll start with a question.

Let X be a finite-dimensional NLS and $Y \subset X$ be a subspace. Given a linear $f : Y \rightarrow \mathbb{R}$, can we extend f to X ? The answer is yes. But what about the infinite-dimensional case? Here, we care about continuous (so bounded) linear operators.

Once again, the answer is that it's possible, but this is hard to prove, and it'll take us a while to prove that. We won't need all the properties of a norm to prove that, so we can weaken what we need in terms of the norm.

Definition. Let X be a vector space over \mathbb{F} . We say that $p : X \rightarrow [0, \infty)$ is *sublinear* if

- (1) $p(\lambda x) = \lambda p(x)$ for all $\lambda \geq 0$ and $x \in X$, and
- (2) $p(x + y) \leq p(x) + p(y)$ for all $x, y \in X$.

If in addition p satisfies (1) for all $\lambda \in \mathbb{F}$, p is called a *seminorm*.

If a seminorm also satisfies $p(x) = 0$ implies $x = 0$, then p is a norm.

The Hahn-Banach theorem about extension of linear operators will apply perfectly well to sublinear operators. First, let's deal with the simplest version we can think of.

Lemma 1.7.4. *Let X be a vector space over \mathbb{R} and $Y \subsetneq X$ be a linear subspace. Let p be sublinear on X and $f : Y \rightarrow \mathbb{R}$ be linear such that $f(y) \leq p(y)$ for all $y \in Y$. For a given $x_0 \in X \setminus Y$, let $\tilde{Y} = \text{span}\{Y, x_0\} = Y + \mathbb{R}x_0 = \{y + \lambda x_0 : y \in Y, \lambda \in \mathbb{R}\}$; then, there exists a linear map $\tilde{f} : \tilde{Y} \rightarrow \mathbb{R}$ such that $\tilde{f}|_Y = f$ and $-p(-x) \leq \tilde{f}(x) \leq p(x)$ for all $x \in \tilde{Y}$.*

The definitions of \tilde{Y} all show that it's "Y plus one more dimension."

PROOF. If $\tilde{f}(x) \leq p(x)$, then $-\tilde{f}(x) = \tilde{f}(-x) \leq p(-x)$, so $\tilde{f}(x) \geq -p(-x)$, and so the lower bound comes for free.

We'll present the proof not as a cleaned-up proof, but how one would think of the proof when trying to prove it.

If we had such an \tilde{f} , what would it look like? $\tilde{y} \in \tilde{Y}$ can be written $\tilde{y} = y + \lambda x_0$ for some $y \in Y$ and $\lambda \in \mathbb{R}$, so $\tilde{f}(\tilde{y}) = \tilde{f}(y + \lambda x_0) = \tilde{f}(y) + \lambda \tilde{f}(x_0) = f(y) + \lambda \tilde{f}(x_0)$, since $\tilde{f}|_Y = f$.

So if we had defined $\alpha \in \mathbb{R}$ to be $\tilde{f}(x_0)$, then we get a function, and correspondingly, given \tilde{f} , we get $\alpha = \tilde{f}(x_0)$. Thus, \tilde{f} is characterized by α .

However, we need to be careful: is this really well-defined? We chose y ; what if you choose a different one than I do? It turns out that you have to choose the same y : suppose $\tilde{y} = y + \lambda x_0 = z + \mu x_0$ for $y, z \in Y$ and $\lambda, \mu \in \mathbb{R}$. Thus, $y - z = (\mu - \lambda)x_0$, but $y - z \in Y$, so since $x_0 \notin Y$, then $\mu - \lambda = 0$, and therefore $y = z$; thus, this choice of y is well-defined, so \tilde{f} really is characterized by α .

So now we need to find an α such that $\tilde{f}(\tilde{y}) = f(y) + \lambda \alpha \leq p(y + \lambda x_0)$. If $\lambda = 0$ this works, so let's focus on $\lambda \neq 0$. Rescale: let $y = -\lambda x$, so we want to show that $f(-\lambda x) + \lambda \cdot \alpha \leq p(\lambda(x_0 - x))$, or $\lambda(-f(x) + \alpha) \leq p(-\lambda(x - x_0))$.

If $\lambda < 0$, then divide by $-\lambda$: $f(x) - \alpha \leq p(x - x_0)$; when $\lambda > 0$, we get a change in sign: $-(f(x) - \alpha) \leq p(-(x - x_0))$. Together, this means $-p(-(x - x_0)) \leq f(x) - \alpha \leq p(x - x_0)$. Rearranging,

$$f(x) - p(x - x_0) \leq \alpha \leq f(x) + p(x_0 - x).$$

This is our requirement; that is, if there's an α that satisfies this for all $x \in Y$, then we have our desired linear functional.

So let $a = \sup_{x \in Y} (f(x) - p(x - x_0))$ and $b = \inf_{x \in Y} f(x) + p(x_0 - x)$. Now we can ignore α and ask, is it true that $a \leq b$? If so, we're done.

Let $x, y \in Y$. Since p is sublinear, then

$$\begin{aligned} f(x) - f(y) &= f(x - y) \leq p(x - y) \\ &\leq p(x - x_0) + p(x_0 - y) \\ \implies f(x) - p(x - x_0) &\leq f(y) + p(x_0 - y). \end{aligned}$$

In the last equation, first take the infimum on the left, which is a , and the right side doesn't change; then, take the supremum on the right, which is b , and the left side doesn't change. Thus $a \leq b$. \square

This proof can be shortened, by starting with α and suddenly magical things happen, but this helps it make more sense and feel more rigorous.

Transfinite Induction and Generalizing Lemma 1.7.4. Applying this inductively, we can extend a finite number of dimensions, and even a countable number of dimensions! However, standard induction doesn't allow us to extend by an uncountable number of dimensions. This will require a technique called transfinite induction, and therefore a brief vacation into set theory.

Definition. A *ordering* on a set \mathcal{S} is a binary relation \preceq such that for all $x, y, z \in \mathcal{S}$,

- (1) $x \preceq x$,
- (2) if $x \preceq y$ and $y \preceq x$, then $x = y$, and

(3) if $x \preceq y$ and $y \preceq z$, then $x \preceq z$.

Not every set can be ordered. However, some can be partially ordered; a *partial order* on a set is the same except that only some pairs $x \preceq y$ are defined, but the same order axioms are satisfied (in particular, $x \preceq x$ is always defined and true, and if $x \preceq y$ and $y \preceq z$, then $x \preceq z$ is defined and true). A *chain* in a partially ordered set \mathcal{S} is a $\mathcal{C} \subset \mathcal{S}$ such that $\preceq|_{\mathcal{C}}$ is a total order: every element can be compared.

Example 1.7.5. On \mathbb{C} , write $z = r_z e^{i\theta_z}$, with $\theta_z \in [0, 2\pi)$.

(1) An ordering on \mathbb{C} can be given by $x \preceq y$ iff $r_x < r_y$ or $r_x = r_y$ and $\theta_x \leq \theta_y$.

(2) A partial ordering on \mathbb{C} can be given by $x \preceq y$ iff $\theta_x = \theta_y$ and $r_x \leq r_y$ (and is undefined if $\theta_x \neq \theta_y$).

We'll need a more complicated order, which requires using Zorn's lemma. This comes from an axiom of set theory called the Axiom of Choice, which states that, given any collection of nonempty sets, it's possible to choose one element out of each set.

Zorn's lemma is equivalent to the Axiom of Choice, but it somehow seems harder to believe.

Lemma 1.7.6 (Zorn's lemma). *Let \mathcal{S} be a nonempty, partially ordered set, and suppose every chain $\mathcal{C} \subseteq \mathcal{S}$ has an upper bound, i.e. for all \mathcal{C} , there's a $u \in \mathcal{S}$ such that $x \preceq u$ for all $x \in \mathcal{C}$. Then, \mathcal{S} has at least one maximal element m , i.e. if $m \preceq x$ for some $x \in \mathcal{S}$, then $x = m$.*

Next time, we'll use this to extend by an uncountable number of dimensions; then, we'll remove the requirement that the base field is real.

Lecture 8.

The Hahn-Banach Theorem, II: 9/14/15

Recall that we're in the middle of proving the Hahn-Banach theorem, and therefore should remember the results we're going to need. We defined orders and partial orders and chains within partially ordered sets last lecture, and cited Zorn's lemma, Lemma 1.7.6, which gives conditions for when a partially ordered set has a maximal element. Finally, we have Corollary 1.7.1 in mind as a long-term goal.

Since we have a possibly countable number of dimensions, we have to use transfinite induction to prove the most general theorem, which is why Zorn's lemma shows up.

Theorem 1.8.1 (Hahn-Banach theorem for real vector spaces). *Let X be a vector space over \mathbb{R} , $Y \subset X$ be a subspace, and p be sublinear on X . If $f : Y \rightarrow \mathbb{R}$ is linear on Y and $f(x) \leq p(x)$ for all $x \in Y$, then there exists a linear $F : X \rightarrow \mathbb{R}$ such that $F|_Y = f$ and $-p(-x) \leq F(x) \leq p(x)$ for all $x \in X$.*

PROOF. Let \mathcal{S} denote the set of all linear extensions g of f to a subspace $D(g) \subset X$ containing Y , and such that $g(x) \leq p(x)$ for all $x \in D(g)$. Since $f \in \mathcal{S}$, then \mathcal{S} is nonempty. We'll turn \mathcal{S} into a partially ordered set by saying that $g \preceq h$ if h extends g , i.e. $D(g) \subseteq D(h)$ and $h|_{D(g)} = g$.

Let \mathcal{C} be a chain in \mathcal{S} , and let

$$D = \bigcup_{g \in \mathcal{C}} D(g).$$

Since these $D(g)$ are nested (i.e. one of $D(g) \subset D(h)$ or $D(g) \supset D(h)$ for all $g, h \in \mathcal{C}$), then D is a vector space.⁸ Then, we'll define $g_{\mathcal{C}}$ as follows: if $x \in D$, then $x \in D(g)$ for some $g \in \mathcal{C}$, so define $g_{\mathcal{C}}(x) = g(x)$. Is this well-defined? Yes, because if $x \in D(g) \cap D(h)$, then without loss of generality $g \preceq h$, and so $g(x) = h(x)$. Thus, we get a function $g_{\mathcal{C}} : D \rightarrow \mathbb{R}$, which is linear (which follows from its definition), and is bounded by p (specifically, $g(x) \leq p(x)$ for all $x \in D$), since each $g \in \mathcal{C}$ is. Thus, $g_{\mathcal{C}} \in \mathcal{C}$, and it's an upper bound for \mathcal{C} .

Applying Zorn's lemma, we have a maximal element F for \mathcal{S} ; since $F \in \mathcal{S}$, then it's a linear extension of f and is bounded by p . So the final question is, what's $D(F)$? Suppose $D(F) \subsetneq X$; then, there exists some $x_0 \in X \setminus D(F)$, so by Lemma 1.7.4 we can extend F to $\text{span}\{D(F), x_0\}$. But this contradicts the fact that $D(F)$ is maximal. Thus, $D(F) = X$. \square

Awesome. Now, let's deal with complex vector spaces. Since we want scalar multiplication for all $\lambda \in \mathbb{C}$, we'll have to use a seminorm instead.

⁸This is an important point; the union of subspaces isn't in general a vector subspace when they're not nested.

Theorem 1.8.2 (Hahn-Banach theorem for complex vector spaces). *Let X be a vector space over \mathbb{F} , $Y \subset X$ be a linear subspace, and p be a seminorm. If $f : Y \rightarrow \mathbb{F}$ is a linear functional such that $|f(x)| \leq p(x)$ for all $x \in Y$, then there exists an extension $F : X \rightarrow \mathbb{F}$ such that $F|_Y = f$ and $|F(x)| \leq p(x)$ for all $x \in X$.*

PROOF. We'll assume $\mathbb{F} = \mathbb{C}$, since the real case comes from Theorem 1.8.1. Then, we can write $f(x) = g(x) + ih(x)$ for g, h real linear, since

$$\begin{aligned} f(x + y) &= g(x + y) + ih(x + y) \\ &= f(x) + f(y) = g(x) + g(y) + ih(x) + ih(y), \end{aligned}$$

and scalar multiplication is similar, though only for real scalars. Instead, $f(ix) = if(x) = -h(x) + ig(x)$, and this is also $g(ix) + ih(ix)$. Thus, $h(x) = -g(ix)$. That is, since f is linear, $f(x) = g(x) - ig(ix)$, which is a general fact.

Since g is real linear, then Theorem 1.8.1 yields a real extension G on X , because $|g(x)| \leq |f(x)| \leq p(x)$, and we have that $|G(x)| \leq p(x)$.

Define $F(x) = G(x) - iG(ix)$, which is a function $F : X \rightarrow \mathbb{C}$ that commutes with addition and real scalar multiplication. Thus, we need to check complex scalar multiplication, and therefore that $F(ix) = iF(x)$. Let's check that:

$$F(ix) = G(ix) - iG(-x) = G(ix) + iG(x) = i(G(x) - iG(ix)).$$

Therefore F is \mathbb{C} -linear. Moreover, if $x \in Y$, then $F(x) = G(x) - iG(ix) = g(x) - ig(ix)$, and therefore $F|_Y = f$ as desired. Thus, the only thing left to check is the bound.

Let $x \in X$, and write $F(x) = re^{i\theta}$. Then,

$$r = |F(x)| = e^{-i\theta} F(x) = F(e^{-i\theta} x) = G(g^{-i\theta}(x)) - iG(-e^{-i\theta} x),$$

but the second term is imaginary, and therefore must be zero. Then,

$$\leq p(e^{-i\theta}(x)) = |e^{-i\theta}|p(x) = p(x). \quad \square$$

As a corollary, notice that $p(x) = \|f\|_{Y^*} \|x\|_X$.

The Hahn-Banach theorem has a great number of corollaries, which provide a lot of insight into NLSes and Banach spaces.

Corollary 1.8.3. *Let X be an NLS and $x_0 \in X \setminus 0$ be fixed. Then, there exists an $f \in X^*$ such that $\|f\|_{X^*} = 1$ and $f(x_0) = \|x_0\|$.*

The idea is to define f on a subspace where it's easy to define, and then extend.

PROOF. Let $Z = \mathbb{F}x_0$, and define $h : Z \rightarrow \mathbb{F}$ by $h(\lambda x_0) = \lambda \|x_0\|$. Then, $|h(x_0)| = |\lambda| \|x_0\| = \|\lambda x_0\|$, so $|h(x)| \leq \|x\|$ for all $x \in Z$ and $\|h\| = 1$. Then, we use Theorem 1.8.2 to extend h to the desired f . \square

Corollary 1.8.4. *For any $\alpha \in \mathbb{F}$, there exists an $f \in X^*$ such that $f(x_0) = \alpha \|x_0\|$ (and therefore $\|f\|_{X^*} = |\alpha|$).*

The proof is the same as for Corollary 1.8.3, but one defines $h(\lambda x_0) = \alpha \lambda \|x_0\|$ instead.

Here's a more interesting corollary.

Corollary 1.8.5. *Let X be an NLS and $x \in X$. Then,*

$$\|x\| = \sup_{\substack{f \in X^* \\ f \neq 0}} \frac{|f(x)|}{\|f\|_{X^*}} = \sup_{\substack{f \in X^* \\ \|f\|_{X^*} = 1}} \frac{|f(x)|}{\|f\|_{X^*}}.$$

Often, when one knows the structure of the dual space better than that of the original space, this can be a useful way to calculate a norm.

PROOF. For all $f \in X^*$ with $f \neq 0$, we know $|f(x)| \leq \|f\|_{X^*} \|x\|$, so we know the supremum is still bounded by $\|x\|$. To get the other bound, we need the Hahn-Banach theorem, which says that there exists a $\tilde{f} \in X^*$ such that $\tilde{f}(x) = \|\tilde{f}\| \|x\|$; then,

$$\sup_{\substack{f \in X^* \\ f \neq 0}} \frac{|f(x)|}{\|f\|_{X^*}} \geq \frac{\tilde{f}(x)}{\|\tilde{f}\|} = \|x\|. \quad \square$$

The idea here is that we can look at $\|x\|$, which is a calculation involving an abstract vector, or $\{f(x)\}_{f \in X^*}$, which is a collection of numbers, which sometimes is nicer. This is a common theme in functional analysis. The following result is related, at least in ideas.

Proposition 1.8.6. *If $f(x) = f(y)$ for all $f \in X^*$, then $x = y$.*

We'll prove this next time.

Lecture 9.

Separability: 9/16/15

Recall that we're in the middle of exploring the consequences of the Hahn-Banach theorem, Theorems 1.8.1 and 1.8.2. For example, if X is an NLS and $x_0 \in X$, then there's an $f \in X^*$ such that $f(x_0) = \|x_0\|$ (Corollary 1.8.4), that you can calculate $\|x\|$ from the norms of $f \in X^*$ (Corollary 1.8.5), and more.

Proposition 1.9.1. *If X is an NLS, then X^* separates points in X , i.e. for any $x, y \in X$, there exists an $f \in X^*$ such that $f(x) \neq f(y)$, and if $f(x) = f(y)$ for all $f \in X^*$, then $x = y$.*

The recurring theme is that if you know what all the linear functionals do to an element, you know what that element is.

PROOF. Choose $x, y \in X^*$ such that $x \neq y$. Then, $x - y = x_0 \in X$ and $x_0 \neq 0$, and there exists an $f \in X^*$ such that $f(x_0) \neq 0$, so $0 \neq f(x - y) = f(x) - f(y)$. \square

Corollary 1.9.2. *If $f(x) = 0$ for all $x \in X^*$, then $x = 0$.*

Oftentimes, one creates simple functionals by doing something interesting on a finite-dimensional subspace and then extending à la the Hahn-Banach theorem.

Definition. In an NLS X , the *distance* between a subspace $Y \subset X$ and a $w \in X$ is $\text{dist}(w, Y) = \inf_{y \in Y} \|w - y\|$.

This is nonnegative, and sometimes it's zero even when $w \notin Y$.

Lemma 1.9.3 (Mazur Separation Thm. I). *Let X be an NLS, $Y \subset X$ be a subspace, and $w \in X \setminus Y$. Suppose $d = \text{dist}(w, Y) > 0$. Then, there exists an $f \in X^*$ with*

- $\|f\| \leq 1$,
- $f(w) = d$, and
- $f(y) = 0$ for all $y \in Y$.

PROOF. Let $Z = Y + \mathbb{F}w$. Then, any $z \in Z$ has a unique representation as $z = y + \lambda w$ for exactly one choice of $y \in Y$ and $\lambda \in \mathbb{F}$ (which we discussed last time).

Then, define $g : Z \rightarrow \mathbb{F}$ by $g(y + \lambda w) = \lambda d$. g is clearly linear, but it's less clear why $\|g\| \leq 1$.

$$\left| g\left(\frac{y + \lambda w}{\|y + \lambda w\|}\right) \right| = \frac{|\lambda|d}{\|y + \lambda w\|} = \frac{d}{\|(1/\lambda)y + w\|}.$$

Since $(1/\lambda)y \in Y$, then $\|(1/\lambda)y + w\| \geq d$, and therefore $d/\|(1/\lambda)y + w\| \leq 1$. Then, we use the Hahn-Banach theorem to extend to X . \square

We'll introduce another notion, entirely topological, which will be useful.

Definition. A topological space X is *separable* if it contains a countable dense subset, i.e. a $\mathcal{D} \subset X$ such that $\overline{\mathcal{D}} = X$.

A space might be large and scary, but if it's separable, then everything is close, and therefore we can get a little control on it.

Example 1.9.4.

- (1) $\mathbb{Q} \subset \mathbb{R}$. \mathbb{Q} is countable and every real number can be arbitrarily well approximated by rational numbers, so \mathbb{R} is separable.
- (2) $\mathbb{Q}(i) = \mathbb{Q} + i\mathbb{Q} \subseteq \mathbb{C}$ is countable and dense, so \mathbb{C} is separable.
- (3) \mathbb{F}^d is also separable, with the countable dense subset either \mathbb{Q}^d or $\mathbb{Q}(i)^d$.
- (4) If $1 \leq p < \infty$, our Schauder basis for ℓ^p is uncountable, but we can take instead the $\mathbb{Q}(i)$ -span (or the \mathbb{Q} -span if $\mathbb{F} = \mathbb{R}$) of $\{e_i\}_{i=1}^\infty$; this is a countable dense subset of ℓ^p , so ℓ^p is separable.

- (5) If $1 \leq p < \infty$, then $L^p(\Omega)$ is separable. This one is a little more surprising. The set S of simple functions (functions which are constant on a finite number of intervals). $S \subset L^p(\Omega)$ is dense, but uncountable, so we have to restrict it in two ways: first, restrict the allowed intervals to have rational coefficients, and then restrict the functions to take on values in $\mathbb{Q}(i)$ (or \mathbb{Q} ; we'll assume that when we talk about $\mathbb{Q}(i)$, then we mean \mathbb{Q} for \mathbb{R}). Thus restricted, we have our countable dense subset.

This argument doesn't work for $L^\infty(\Omega)$, since simple functions aren't dense in it, and in fact $L^\infty(\Omega)$ isn't separable.

Proposition 1.9.5. *Let X be an NLS. If X^* is separable, then X is separable.*

The converse isn't true, because $L^1(\Omega)$ is separable but $L^\infty(\Omega)$ isn't. So if you start with a separable space, your dual might be bigger.

PROOF. Let $\{f_n\}_{n=1}^\infty$ be a countable, dense subset of X^* . We'll use this to construct a countable, dense subset of X . Since $\|f\| = \sup_{\|x\|=1} |f(x)|$, then we can choose for each n an x_n such that $\|x_n\| = 1$ and $|f_n(x_n)| \geq (1/2)\|f_n\|$, giving us a sequence $\{x_n\}_{n=1}^\infty$.

Then, let $\mathcal{D} = \text{span}_{\mathbb{Q}(i)}\{x_n\}$, which is still countable, and we'll show that $\overline{\mathcal{D}} = X$. Suppose that it weren't: then, there exists a $w \in X \setminus \overline{\mathcal{D}}$. Let $d = \text{dist}(w, \mathcal{D}) = \inf_{x \in \mathcal{D}} \|w - x\| > 0$. If we can show that $\|w - y_n\| \rightarrow 0$ for some sequence $\{y_n\}_{n=1}^\infty$, then since $\overline{\mathcal{D}}$ is closed, that would imply $w \in \mathcal{D}$.

Since $\overline{\mathbb{Q}(i)} = \mathbb{C}$, then $\overline{\mathcal{D}}$ is a linear subspace of X ; thus, by Lemma 1.9.3, there exists an $f \in X^*$ such that $f|_{\overline{\mathcal{D}}} = 0$ and $f(w) = d > 0$. But there's a sequence $\{f_{n_k}\}_{k=1}^\infty$ such that $f_{n_k} \rightarrow f$. Thus,

$$\|f_{n_k} - f\| \geq |f(x_{n_k}) - f_{n_k}(x_{n_k})| = |f_{n_k}(x_{n_k})| \geq \frac{1}{2}\|f_{n_k}\|.$$

Since $f_{n_k} - f \rightarrow 0$, then this means $f_{n_k} \rightarrow 0$, and so $f = 0$. But this is a contradiction. \square

So far, we've always looked at sets that are subspaces. Here's an example where we don't do that.

Definition. Let X be an NLS and $C \subseteq X$ be a subset (not necessarily a subspace). Then, C is *balanced* if for any $\lambda \in \mathbb{F}$ with $|\lambda| \leq 1$, then when $x \in C$, $\lambda x \in C$.

For example, if $\mathbb{F} = \mathbb{C}$, then this implies that C is invariant under rotation, as well as contractions. Note that all subspaces are balanced.

Lemma 1.9.6 (Mazur Separation Thm. II). *Let X be an NLS and $C \subseteq X$ be a closed, convex, and balanced set. Then, for any $w \in X \setminus C$, there exists an $f \in X^*$ such that $|f(x)| \leq 1$ for $x \in C$ and $f(w) > 1$.*

PROOF. Since C is closed and $w \notin C$, we can choose a ball $B + w$ about w (so B is a ball centered at the origin) such that $B \cap C = \emptyset$. Then, we can define the *Minkowski functional* $p : X \rightarrow [0, \infty)$ by

$$p(x) = \inf\left\{t > 0 : \frac{x}{t} \in C + B\right\}.$$

Here, $C + B$ is a slight fattening of our set C , but we can guarantee that $w \notin C + B$. Moreover, $0 \in C$, because C is balanced; therefore, $p(x)$ is always finite. We also know that $p(x) \leq 1$ if $x \in C$ and $p(w) > 1$.

Moreover, p is a seminorm: since C is balanced, $p(\lambda x) = p(|\lambda|x) = |\lambda|p(x)$. We also have the triangle inequality, which is left to the reader.

Now, we use Theorem 1.8.2: let $Y = \mathbb{F}w$, and if $f(\lambda w) = \lambda p(w)$, then $f(w) = p(w) > 1$, and $|f(\lambda w)| = |\lambda|p(w) = p(\lambda w)$, so we have a nice bound. Therefore, we can extend f to an F such that $F(w) > 1$ and $F(x) \leq 1$ if $x \in C \subset C + B$. \square

Basically, the idea is the Minkowski functional; once you write that down, you're basically done.

Lecture 10.

The Minkowski Functional and the Baire Category Theorem: 9/18/15

Last time, we had to rush through the Minkowski functional, so today we'll talk a little more about it. This is *not* a linear functional, but it does map into \mathbb{F} , so it's called a functional.

Specifically, given a nonempty $A \subseteq X$, where X is an NLS, the Minkowski functional is defined as

$$p(x) = \inf\{t > 0 : x \in tA\},$$

which takes values in $[0, \infty]$. We then showed the following.

- (1) If there's an open ball containing 0 and contained in A , then $p(x)$ is finite.
- (2) p is *positively homogeneous*, i.e. if $\lambda \geq 0$, then $p(\lambda x) = \lambda p(x)$.
- (3) If A is convex, then $p(x + y) \leq p(x) + p(y)$. Well, we didn't actually show (3), so let's do that now. Suppose $x/r, y/s \in A$ (so that $r \geq p(x)$ and $s \geq p(y)$). By convexity,

$$\frac{x+y}{s+r} = \frac{r}{s+r} \frac{x}{r} + \frac{s}{s+r} \frac{y}{s} \in A,$$

and therefore $s+r \geq p(x+y)$. Since this is true for all such s and r , passing to their infimum replaces them with $p(x)$ and $p(y)$, so $p(x) + p(y) \geq p(x+y)$.

- (4) We did show that if A is balanced, then p is a seminorm.

This was sufficient to prove Lemma 1.9.6, but we have one more separating theorem to prove. This time, we don't need sets to be balanced, but we will require convexity.

Lemma 1.10.1 (Separating hyperplane theorem). *Let A and B be disjoint, nonempty, convex subsets of an NLS X .*

- (1) *If A is open, then there exists an $f \in X^*$ and a $\gamma \in \mathbb{R}$ such that $\operatorname{Re}(f(x)) \leq \gamma \leq \operatorname{Re}(f(y))$ for all $x \in A$ and $y \in B$.*
- (2) *If A and B are open, the above inequality is strict.*
- (3) *If A is compact and B is closed, then the above inequality is also strict.*

PROOF. We'll prove part 1; the others are similar. Moreover, it suffices to prove it for real fields, because if $\mathbb{F} = \mathbb{C}$, then we can view X as a real vector space and get a real linear functional g that satisfies the lemma over \mathbb{R} . Then, $f(x) = g(x) - ig(ix)$ satisfies the lemma for \mathbb{C} .

All right, so $\mathbb{F} = \mathbb{R}$, and A is open and both are convex. We'll have to put the Minkowski functional into this proof somehow, so let's start by picking an $x \in A$ and a $y \in B$. Let $A-x = \{t-x : t \in A\}$, and define $B-y$ similarly. Then, let $C = (A-x) - (B-y) = \{t-s-x+y : t \in A, s \in B\}$, and for convenience, let $w = y-x$. We'll want to construct a Minkowski functional on C .

C is open, since A is; convex, because A and B are; and contains 0 (because we've moved x and y to the origin). But $w \notin C$, since A and B are disjoint. Let $Y = \mathbb{R}w$ and $g(tw) = t$; then, our Minkowski functional is $p(x) = \inf\{t > 0 : x \in tC\}$, which is well-defined and sublinear, and satisfies $p(w) \geq 1 = g(w)$, so $g(w) \leq p(w)$, and therefore for any $y \in Y$, $g(y) \leq p(y)$: if $\lambda \geq 0$, this follows from the positive homogeneity of p and the linearity of g , and if $\lambda < 0$, $-\lambda g(w) \geq -\lambda p(w)$, and therefore $-g(-\lambda w) \leq p(-\lambda w)$.

Thus, we can extend g to X , and $g(x) \leq 1$ on C , since $g(x) \leq p(x)$ everywhere on X , and therefore $g(x) \geq -1$ for $x \in -C$, and so $|g(x)| \leq 1$ on $C \cap (-C)$. Since this contains a neighborhood of the origin, g is bounded, so $g \in X^*$.

If $a \in A$ and $b \in B$, then $a-b+w \in C$, and therefore $1 \geq g(a-b+w) = g(a) - g(b) + 1$, so $g(b) \geq g(a)$. Let $\gamma = \sup g(a)$ (or $\gamma = \inf g(b)$), and we're done. \square

This concludes our discussion of the Hahn-Banach theorem and its applications.

The Open Mapping Theorem. The open mapping theorem, which is the next major result for Banach spaces, helps us characterize what linear functionals can look like.

The following theorem is important in its own right, but we'll use it as an ingredient in the proof of the open mapping theorem.

Theorem 1.10.2 (Baire category theorem). *Let (X, d) be a complete metric space, and let $\{V_j\}_{j=1}^\infty$ be a sequence of open, dense subsets of X . Then, $V = \bigcap_{j=1}^\infty V_j$ is dense.*

In other words, open dense sets aren't exactly thin: they're actually surprisingly fat, so fat that a countable intersection of them is still fat, in a sense.

PROOF. Let $W \subseteq X$ be any nonempty open set. Then, we have to show that $V \cap W \neq \emptyset$, since V being dense is equivalent to intersecting every nonempty open.

Since V_1 is dense, then $W \cap V_1 \neq \emptyset$, so there's an $x_1 \in W \cap V_1$, and since W is open, there's a $\overline{B_{r_1}(x_1)} \subseteq W \cap V_1$ (we have an open neighborhood, and can take the closure of a smaller ball); also, we can without loss of generality take $0 < r_1 < 1$, by shrinking r_1 if necessary.

In the same way, V_2 is open and dense and $B_{r_1}(x_1)$ is a nonempty open set, so there exists an $x_2 \in V_2 \cap B_{r_1}(x_1)$ and an $r_2 \in (0, 1/2)$, such that $\overline{B_{r_2}(x_2)} \subseteq B_{r_1}(x_1) \cap V_2$. Then, we can continue in this way, choosing for each n an x_n and an r_n such that $\overline{B_{r_n}(x_n)} \subseteq B_{r_{n-1}}(x_{n-1}) \cap V_n$ and $0 < r_n < 1/n$.

Now, consider the sequence $\{x_n\}_{n=1}^\infty$, which is Cauchy, because if $i, j \geq n$, then $x_i, x_j \in B_{r_n}(x_n)$ and therefore $d(x_i, x_j) < 2/n$. Since X is complete, this sequence converges to some $x \in X$. Since $x_i \in B_{r_n}(x_n)$ for all $i > n$, then $x \in \overline{B_{r_n}(x_n)}$, so $x \in V_n$ for all n . Since $x \in \overline{B_{r_1}(x_1)} \subseteq W$, then $x \in W \cap V$, so the intersection is nonempty, and thus V is dense. \square

Some of you may have been disappointed to see that no category theory appeared in the statement or proof; in this part of mathematics, “category” has a different definition.

Definition. Let (X, d) be a metric space.

- A is *nowhere dense* if it has empty interior: $(\overline{A})^0 = \emptyset$.
- A is *first category* if it can be written as a countable union of nowhere dense sets.
- If A isn't first category, then it is called *second category*.

Using these definitions, the Baire category theorem says that a complete metric space is second category.

Corollary 1.10.3. *A complete metric space is not the countable union of nowhere dense sets.*

In other words, a complete metric space is fatter than that.

PROOF. Suppose X is such a union: $X = \bigcup_{j=1}^\infty M_j$ with each M_j nowhere dense. Without loss of generality, each M_j is closed (or just take their closures, which still cover X). Thus, by de Morgan's law, $\emptyset = \bigcap_{j=1}^\infty (M_j)^c$. Since M_j is closed and nowhere dense, then M_j^c is a dense open set, and therefore \emptyset is the countable intersection of dense open sets, which contradicts Theorem 1.10.2. \square

Next time, we'll return to the world of Banach spaces.

Lecture 11.

The Open Mapping Theorem: 9/21/15

Last time, we learned about the Baire category theorem. Today, we'll use it to prove the open mapping theorem.

Definition. A continuous map $f : X \rightarrow Y$ is *open* if it maps open sets to open sets, i.e. if $U \subseteq X$ is open, then $f(U) \subseteq Y$ is open.

An arbitrary continuous map is not open; for example, $T : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ sending $(x, y) \mapsto (x, 0)$ is perfectly continuous, but the image of $B_1(0)$ is $(0, 1) \times \{0\}$, which isn't open in \mathbb{R}^2 .

In the infinite-dimensional case, things can become more interesting; for example, $T : \ell^2 \rightarrow \ell^2$ sending $e_n \mapsto (1/n)e_n$ isn't open (the image of the unit ball isn't open), but is linear and surjective; the discrepancy is that this T isn't bounded.

Theorem 1.11.1 (Open mapping). *Let X and Y be Banach and $T : X \rightarrow Y$ be a bounded, linear surjection. Then, T is an open map.*

A bounded linear map is typical in this class; the key hypothesis in this theorem is that T is surjective; the example $(x, y) \mapsto (x, 0)$ shows that this is important.

This is a pretty fundamental theorem about Banach spaces.

PROOF. It suffices to show that $T(B_1(0))$ contains a $B_r(0)$ for some $r > 0$: if $U \subset X$ is open, then to check that $T(U)$ is open, we can pick a $y \in T(U)$ and a preimage x (i.e. $T(x) = y$). Then, we can look at $U - x$, and since T is linear, then $T(U - x) = T(U) - y$. But since x and y are now sent to the origin, we just need to pick a neighborhood of x and make sure its image contains a neighborhood of y .

This is the proper way to think about the theorem: if you know what a bounded linear map looks like at the origin, you know what it looks like everywhere.

Since T is onto, then we can write

$$Y = \bigcup_{j=1}^{\infty} T(B_j(0)).$$

Since Y is a complete metric space, then the Baire category theorem tells us it's not the union of nowhere dense sets. Thus, there's some k such that $T(B_k(0))$ isn't nowhere dense, i.e. there's an open $W_1 \subset \overline{T(B_k(0))}$. Thus, we can scale: $(1/2k)W \subseteq (1/2k)\overline{T(B_k(0))} = \overline{T(B_{1/2}(0))}$.

Since W_1 is open, there's a $y_0 \in Y$ and an $r > 0$ such that $B_r(y_0) \subseteq W \subseteq \overline{T(B_{1/2}(0))}$. This is almost everything:

$$\begin{aligned} B_r(0) &= B_r(y_0) - y_0 \\ &\subseteq B_r(y_0) - B_r(y_0) \\ &\subseteq \overline{T(B_{1/2}(0))} + \overline{T(B_{1/2}(0))} \\ &\subseteq \overline{T(B_1(0))}, \end{aligned}$$

by the triangle inequality. We'd be done, except that we had to take the closure (which ultimately came from the Baire category theorem). Thus, we'll show that if $\varepsilon > 0$, then $T(B_{1+\varepsilon}(0)) \supseteq B_r(0)$, because then

$$T(B_1(0)) = \frac{1}{1+\varepsilon} T(B_{1+\varepsilon}(0)) \supseteq B_{r/(1+\varepsilon)}(0).$$

Then, we won't need the closure anymore. Note that this isn't obvious, even if it seems obvious in the finite-dimensional case.

Fix a $y \in B_r(0)$ and an $\varepsilon > 0$. We know that $T(B_1(0)) \cap B_r(0)$ is dense in $B_r(0)$ (since we showed already its closure contains $B_r(0)$), so we can pick an $x_1 \in B_1(0)$ such that $\|y - Tx_1\| \leq \varepsilon/2$.

Inductively, when $n \geq 1$, suppose we've picked x_1, \dots, x_n such that $\|x_1\| \leq 1$ and $\|x_j\| \leq 2^{-j+1}\varepsilon$ and $\|y - T(x_1 + \dots + x_n)\| < 2^{-n}\varepsilon r$. Let $z = y - T(x_1 + \dots + x_n)$, so that $z \in B_{2^{-n}\varepsilon r}(0)$. Since $T(B_1(0)) \cap B_r(0)$ is dense in $B_r(0)$, we can scale things: there's an $x_{n+1} \in B_{2^{-n}\varepsilon}(0)$ such that $\|z - Tx_{n+1}\| \leq 2^{-(n+1)}\varepsilon r$; thus, $\|y - T(x_1 + \dots + x_{n+1})\| \leq 2^{-(n+1)}\varepsilon r$.

Since the terms get smaller and smaller, $\sum_{j=1}^n x_j$ is a Cauchy sequence, so since X is complete, then this sum converges to a point $x \in X$, such that

$$\|x\| \leq 1 = \sum_{j=2}^{\infty} \|x_j\| < 1 + \sum_{n=2}^{\infty} 2^{-n+1}\varepsilon = 1 + \varepsilon,$$

and T is continuous, then $Ts_n \rightarrow Tx = y$. ⊠

The first part, showing it's true for $\overline{T(B_1(0))}$, is pretty easy, but then getting just one more ε is surprisingly fussy.

Corollary 1.11.2. *If X and Y are Banach spaces and $T : X \rightarrow Y$ is a bounded, linear bijection (one-to-one and onto), then the inverse map exists and is a bounded linear functional.*

In other words, the inverse of a bounded linear functional is bounded linear. This is nice, and very useful.

PROOF. It's easy to show T^{-1} is linear. T is open, so it takes open sets to open sets, and therefore for T^{-1} , the preimage of every open set is open, so T^{-1} is continuous. ⊠

We now know enough to make the following definition.

Definition. If X and Y are Banach spaces, we say that they're *isomorphic as Banach spaces* if there exists a linear, bounded bijection $T : X \rightarrow Y$. If in addition T preserves the norm, it's called a *isometry*.

This means that X and Y have the same vector space structure (since there's a bijective linear map) and same topology (there's a homeomorphism). If T isn't an isometry, then the norms may be different, but they'll be equivalent, so X and Y are basically the same.

There's a closely related result about graphs of maps; we could have proven this first and used it to derive the open mapping theorem, though we'll go about it in the other direction.

Definition. Let X and Y be topological spaces, $D \subseteq X$ and $f : D \rightarrow Y$. Then, the *graph* of f is $\text{graph}(f) = \{(x, f(x)) : x \in D\} \subseteq X \times Y$.

$X \times Y$ is where we're used to drawing graphs (such as $X, Y = \mathbb{R}$); we chose D because the function might not be defined everywhere.

Proposition 1.11.3. *Let X be a topological space, Y be a Hausdorff space, and $f : X \rightarrow Y$ be continuous. Then, $\text{graph}(f)$ is closed in $X \times Y$.*

In the case of graphs we're most familiar with, this makes sense, as it's how we're used to thinking of continuity intuitively.

PROOF. Let $U = X \times Y \setminus \text{graph}(f)$, and let $(x_0, y_0) \in U$, so $f(x_0) \neq y_0$.

Recall that since Y is Hausdorff, we can choose open neighborhoods $V, W \subset Y$ around y_0 and $f(x_0)$, respectively, that don't intersect. Then, we're going to consider $f^{-1}(W) \subset X$; specifically, $f^{-1}(W) \times V$ doesn't intersect $\text{graph}(f)$, and is an open neighborhood of (x_0, y_0) . \square

If that proof didn't make sense, drawing a picture will likely help.

Definition. Let X and Y be NLSes, $D \subseteq X$ be a linear subspace, and $T : D \rightarrow Y$ be linear. Then, we say that T is a *closed operator* if $\text{graph}(T)$ is closed in $X \times Y$.

An open map takes open sets to open sets; a bounded map takes bounded sets to bounded sets; but a closed operator *doesn't* take closed sets to closed sets. This can be confusing.

If X and Y are metric spaces and $f : D \rightarrow Y$ is continuous, then $\text{graph}(f)$ being closed means that if $\{x_n\}_{n=1}^\infty \subseteq D$ such that $x_n \rightarrow x$ and $Tx_n \rightarrow y$, then $x \in D$ and $y = Tx$.


Theorem 1.11.4 (Closed graph theorem). *Let X and Y be Banach spaces, and $T : X \rightarrow Y$ be linear. Then, T is bounded iff it's closed.*

PROOF. The forward direction is true in general (continuous implies closed).

In the other direction, suppose $\text{graph}(T)$ is closed, and therefore is a closed linear subspace of $X \times Y$. This is a very important point: since X and Y are Banach, then $X \times Y$ is Banach, and since $\text{graph}(T)$ is closed in it, then $\text{graph}(T)$ is also a Banach space, with the *graph norm*⁹

$$\|(x, Tx)\| = \|x\|_X + \|Tx\|_Y.$$

Define two projection operators $\pi_1 : (x, y) \mapsto x$ and $\pi_2 : (x, y) \mapsto y$, so that we have maps



π_1 is a linear bijection between Banach spaces, and is bounded (by the triangle inequality, $\|x\|_X \leq \|(x, Tx)\| + \|Tx\|_Y$), so by Corollary 1.11.2, its inverse π_1^{-1} is a bounded linear functional. Moreover, π_2 is bounded linear for the same reasons, so $T = \pi_2 \circ \pi_1^{-1}$ is bounded as well. \square

Lecture 12.

The Uniform Boundedness Principle: 9/23/15

Recall that last time, we proved the closed graph theorem, Theorem 1.11.4, which states that if X and Y are Banach spaces and $T : X \rightarrow Y$ is linear, then T has a closed graph iff it's bounded.

Corollary 1.12.1. *Let X and Y be Banach spaces, $D \subset X$ be a subspace, and $T : D \rightarrow Y$ be closed. Then, T is bounded iff D is closed.*

PROOF. In the reverse direction, if D is closed, then it's Banach, so we can apply Theorem 1.11.4.

Conversely, suppose T is bounded, and let $\{x_n\}_{n=1}^\infty \subseteq D$ such that $x_n \rightarrow x$ in X . Since T is bounded and linear, then $\{Tx_n\}_{n=1}^\infty$ is Cauchy.¹⁰

⁹Though this was defined in a way mirroring the ℓ^1 norm, you can use the analogous definition with any ℓ^p including ℓ^∞ , since, as we proved on a problem set, they're all equivalent.

¹⁰This is because $\|Tx_n - Tx_m\| = \|T(x_n - x_m)\| = \|T\| \|x_n - x_m\| \rightarrow 0$.

Since Y is complete, then $Tx_n \rightarrow y$ for some $y \in Y$. And since T is a closed operator, then $\text{graph}(T) = \{(x, Tx)\}$ is closed. Since $x_n \rightarrow x$ and $Tx_n \rightarrow y$, then $(x, y) \in \text{graph}(T)$, so $y = Tx$ and thus $x \in D$. Therefore D contains its limit points, and so is closed. \square

Hopefully this illustrates some uses of the closed graph theorem.

Example 1.12.2. Though continuous implies closed, the converse isn't true. Here's an example. Let $X = C([0, 1])$, the continuous functions with the L^∞ norm, and let $D = C^1([0, 1]) \subset X$, the C^1 functions. Let $T : D \rightarrow X$ be the derivative operator, so $T(f) = f'$.

This is perfectly well defined: if f is C^1 , then f' is continuous. Then, $D \neq X$ (e.g. $f(x) = |x - 1/2|$), but $\overline{D} = X$, so D is not closed in X (intuitively, any continuous but not differentiable function can be well approximated by a C^1 function).

- First, we'll see that T isn't continuous (equiv. bounded). $T(x^n) = nx^{n-1}$, but $\|x_n\| = 1$ and $\|Tx_n\| = n \rightarrow \infty$.
- However, it is closed. Let $\{f_n\}_{n=1}^\infty \subseteq D$ have a limit $f_n \rightarrow f$ in X , and such that $f'_n \rightarrow g$ in X . We want to show that $g = f'$. This follows from the fundamental theorem of calculus: for each n ,

$$f_n(t) = f_0(t) + \int_0^t f'_n(\tau) d\tau.$$

Since convergence in L^∞ implies pointwise convergence. Then, by the dominated convergence theorem, these integrals also converge (recall that continuous functions on compact sets are bounded), so

$$f(t) = f(0) + \int_0^t g(\tau) d\tau = f(0) + \int_0^t f'(\tau) d\tau.$$

Thus, $f'(t)$ exists and $g = f'$, so $f \in C^1([0, 1])$.

Continuous is better than closed, but closed is part of the way there, in some sense.

We've talked about two of the three important theorems about NLSes: the Hahn-Banach theorem and the open mapping theorem. Here's the third.

Theorem 1.12.3 (Uniform boundedness principle). *Let X be Banach, Y be an NLS, and $\{T_\alpha\}_{\alpha \in I} \subseteq B(X, Y)$. Then, one of the following holds.*

- (1) *The collection is uniformly bounded: there's an M such that $\sup_{\alpha \in I} \|T_\alpha\|_{B(X, Y)} \leq M$,*
- (2) *There's a point where they're not: there's an $x \in X$ such that $\sup_{\alpha \in I} \|T_\alpha x\| = \infty$.*

The point is, if these functions aren't uniformly bounded, then they all blow up at a given point, when *a priori* they could do so in different places. This is true no matter how large the collection I is; it could very well be uncountable.

The proof in the notes is nice, but a little fussy to prove, and uses the Baire category theorem. We'll give a proof based on the following lemma, which is a little nicer.

Lemma 1.12.4. *Let X and Y be NLSes and $T : X \rightarrow Y$ be a bounded linear map. For any $x \in X$ and $r > 0$, $\sup_{y \in B_r(x)} \|Ty\| \geq r\|T\|$.*

The idea is if $x = 0$, we have equality, but even if we don't, this is still a one-sided bound.

PROOF. For a visualization, it may help to think about the case $X = Y = \mathbb{R}$, where $\text{graph}(T)$ is a line with slope $\|T\|$ through the origin. Here, $\sup_{y \in B_r(x)} \|Ty\| = \|T\|(\|x\| + r) \geq \|T\|r$.

More generally, we'll think of the "larger" and "smaller" parts of $B_r(x)$. The triangle inequality tells us that since $z = (1/2)(x + z) - (x - z)$, then

$$\begin{aligned} \|Tz\| &\leq \frac{1}{2}(\|T(x + z)\| + \|T(x - z)\|) \\ &\leq \max\{\|T(x + z)\|, \|T(x - z)\|\}. \end{aligned}$$

If we take the supremum over $z \in B_r(0)$, we know that

$$r\|T\| \leq \sup_{z \in B_r(0)} \|T(x + z)\| = \sup_{y \in B_r(x)} \|Ty\|.$$

\square

This is a nice geometric result, and relatively easy to prove. Then, we'll use it to attack the uniform boundedness principle.

PROOF OF THEOREM 1.12.3. Since we want to show one of two things is true, let's assume $\sup_\alpha \|T_\alpha\| = \infty$. Choose a countable subcollection $\{T_n\}_{n=1}^\infty$ such that $\|T_n\| \geq 4^n$. Then, choose $x_0 = 0 \in X$ and choose $x_n \in X$ such that $\|x_n - x_{n-1}\| \leq 3^{-n}$ (so that $x_n \in B_{3^{-n}}(x_{n-1})$), so that by Lemma 1.12.4, $\|T_n x_n\| \geq (2/3)3^{-n}\|T_n\|$.

Since $\{x_n\}$ is Cauchy, then it converges to some $x \in X$, and $\|x - x_n\| \leq (1/2)3^{-n}$, since

$$\begin{aligned} \|x - x_n\| &= \lim_{m \rightarrow \infty} \|x_m - x_n\| \\ &\leq \lim_{m \rightarrow \infty} \sum_{j=m}^{n+1} \|x_j - x_{j-1}\| \\ &\leq \lim_{m \rightarrow \infty} (3^{-m} + 3^{-(m-1)} + \dots + 3^{-(n+1)}) \\ &\leq \lim_{m \rightarrow \infty} 3^{-n} \sum_{k=1}^{\infty} \left(\frac{1}{3}\right)^k = \frac{1}{2}3^{-n}, \end{aligned}$$

since it's a nice old geometric series.

Finally, we're going to look at $T_n x$.

$$\|T_n x\| \leq \|T_n(x - x_n)\| + \|T_n x_n\| \leq \|T_n\| \|x - x_n\| + \|T_n x_n\|.$$

Thus, we know that

$$\frac{2}{3}3^{-n}\|T_n\| \leq \|T_n\| \frac{1}{2}3^{-n} + \|T_n x_n\|,$$

and therefore

$$\|T_n x_n\| \geq \frac{1}{6}3^{-n}\|T_n\| \geq \frac{1}{3}\left(\frac{4}{3}\right)^n,$$

which goes to infinity. □

The Double-Dual. If X is an NLS, then $X^* = B(X, \mathbb{F})$ is a Banach space, and we can form the *double-dual* $X^{**} = B(X^*, \mathbb{F})$, which is also a Banach space. It's possible to interpret X as sitting inside X^{**} .

For any $x \in X$, define the evaluation map $E_x : X^* \rightarrow \mathbb{F}$ by $E_x(f) = f(x)$: that is, we evaluate f at x . Since $f : X \rightarrow \mathbb{F}$, then E_x is well-defined, and it's linear: if $f, g \in X^*$ and $\lambda \in \mathbb{F}$, then

$$\begin{aligned} E_x(f + g) &= (f + g)(x) = f(x) + g(x) = E_x(f) + E_x(g) \\ E_x(\lambda f) &= (\lambda f)(x) = \lambda(f(x)) = \lambda E_x(f). \end{aligned}$$

E_x is also bounded, which is more interesting. If you think about what the norm means, then

$$\|E_x\|_{X^{**}} = \sup_{\substack{f \in X^* \\ f \neq 0}} \frac{|E_x(f)|}{\|f\|_{X^*}} = \sup_{\substack{f \in X^* \\ f \neq 0}} \frac{f(x)}{\|f\|_{X^*}} = \|x\|_X,$$

where the last equality is due to Corollary 1.8.5.

Definition. Let (M, d) and (N, ρ) be metric spaces. Then, $f : M \rightarrow N$ is an *isometry* if $\rho(f(x), f(y)) = d(x, y)$ for all $x, y \in M$. If f is surjective, M and N are said to be *isometric*.

Note that f is always injective, because if $f(x) = f(y)$, then $\rho(f(x), f(y)) = 0 = d(x, y)$, so $x = y$. Thus, isometric spaces are given by a bijection f .

Anyways, that's what's going on here: not only is the metric the same, but the norm is the same. We have a map $E : X \rightarrow X^{**}$ sending $x \mapsto E_x$. E is a bounded linear map, and an isometry. Therefore, $\tilde{X} = \{E_x \in X^{**} : x \in X\} \subset X^{**}$ is isomorphic and isometric to X . It might not be all of X^{**} , but we've embedded X into its double-dual.

Definition. Sometimes, $X = X^{**}$ (i.e. $\tilde{X} = X^{**}$). If this is true, X is said to be *reflexive*.

In general, $X \subseteq X^{**} \subseteq X^{****} \subseteq \dots$; if $X = X^{**}$, then this entire chain collapses to equalities. Similarly, we could have started with X^* and X^{***} , and so on.

Theorem 1.12.5. If X is reflexive, then X^* is reflexive.

This is left to the exercises, but isn't hard to prove. Notice, however, that the converse isn't true.

Example 1.12.6. If $1 \leq p < \infty$, we know that $(\ell^p)^* = \ell^q$, so ℓ^p is reflexive for $1 < p < \infty$. ℓ^1 and ℓ^∞ aren't as nice, so we don't have reflexivity. (The duality follows from something on the homework this week.)

Since $(L^p(\Omega))^* = L^q(\Omega)$ for $1 < p < \infty$, then $L^p(\Omega)$ is reflexive (which follows from the Radon-Nikodym theorem). Similarly, $L^1(\Omega)$ and $L^\infty(\Omega)$ are more complicated.

Lecture 13.

Weak and Weak-* Convergence: 9/25/15

In a finite-dimensional NLS, we know that we're basically looking at \mathbb{R}^d , where we have the Heine-Borel theorem: a set is closed and bounded iff it's compact. But, of course, infinite dimensions are weirder, and in fact the unit ball isn't compact in an infinite-dimensional vector space. This was the reasoning behind one of the less intuitive HW problems, about embedding infinitely many disjoint balls of a fixed radius into the unit ball.

Theorem 1.13.1. *Let $Y \subseteq X$ be a closed subspace, and Z be another subspace containing Y . If $Z \neq Y$ and $\theta \in (0, 1)$, then there exists a $z \in Z$ such that $\|z\| = 1$ and $\text{dist}(z, Y) = \theta$.*

The intuition in the finite-dimensional case is that z should be “orthogonal” to Y , but not all infinite-dimensional spaces have an inner product structure (which is what gives us angles), so we have to be careful.

PROOF. Pick a $z_0 \in Z \setminus Y$ and let

$$d = \text{dist}(z_0, Y) = \inf_{y \in Y} \|y - z_0\|.$$

Since Y is closed, $d > 0$, so we can choose a $y_0 \in Y$ such that $d/\theta \geq \|z_0 - y_0\| \geq d$. Then, let $z = (z_0 - y_0)/\|z_0 - y_0\| \in Z$. For all $y \in Y$,

$$\begin{aligned} \|z - y\| &= \frac{\|z_0 - y_0 - y\| \|z_0 - y_0\|}{\|z_0 - y_0\|} \\ &= \frac{\|z_0 - y_1\|}{\|z_0 - y_0\|} \\ &\geq \|z_0 - y_1\| \frac{\theta}{d} \geq \theta. \end{aligned}$$

⊠

Corollary 1.13.2. *If X is an infinite-dimensional NLS and $M \subset X$ is a closed, bounded set with nonempty interior, then M is not compact.*

The Heine-Borel theorem is false in infinite dimensions. Alas.

PROOF. It's sufficient to show it for the (closed) unit ball, because then translations and scalings cover all such closed, bounded sets.

Let $x_1 \in X$ with $\|x_1\| = 1$, and we'll inductively assume we've chosen x_1, \dots, x_n such that $\|x_i\| = 1$ and $\|x_i - x_j\| \geq 1/2$ for all $i \neq j$. Let $Y = \text{span}\{x_1, \dots, x_n\}$, so since Y is finite-dimensional and X is infinite-dimensional, we can choose an $x \in X \setminus Y$ and let $Z = \text{span}\{x, Y\}$. With $\theta = 1/2$, Theorem 1.13.1 gives us an $x_{n+1} \in Z$ with $\|x_{n+1}\| = 1$ and $\text{dist}(x_{n+1}, Y) \geq 1/2$.

Thus, we can pick an infinite sequence of points on the unit sphere such that all of them are at least distance $1/2$ from each other. Thus, $\{x_n\}$ does not converge, so $\overline{B_1(0)}$ cannot be compact.¹¹ ⊠

A modification of this proof shows that the volume is arbitrarily large too, not just the surface area.

What we'll do about this is to define a new topology, which is weaker (has fewer open sets); since compactness is a condition on open covers, this makes things more likely to be compact. We'll start with a new notion of convergence, and then extract the topology afterwards.

Definition. Let X be an NLS and $\{x_n\}_{n=1}^\infty \subset X$.

- We say that x_n converges weakly to x , written $x_n \rightarrow x$ or $x_n \xrightarrow{w} x$, if for any $f \in X^*$, $f(x_n) \rightarrow f(x)$.

¹¹Here, we're using the theorem that compactness is equivalent to sequential compactness in metric spaces, and we should also remark that no subsequence of $\{x_n\}$ converges, which is the criterion we need.

- If $\{f_n\}_{n=1}^\infty$, then f_n converges weak-* (said “weak-star”) to f , written $f_n \xrightarrow{w^*} f$, if for all $x \in X$, $f(x_n) \rightarrow f(x)$.

As a contrast, $x_n \rightarrow x$ given by $\|x_n - x\| \rightarrow 0$ is sometimes called *strong convergence*.

So in a space X^* that’s the dual of some X , we have three notions of convergence floating around: the strong topology $\|f_n - f\|_{X^*} \rightarrow 0$, the weak topology $F(f_n) \rightarrow F(f)$ for all $F \in X^{**}$, and weak-*, where we only consider evaluation maps in X^{**} .

These notions of convergence represent new topologies. Let’s establish some propositions before we continue.

Proposition 1.13.3. *If $x_n \rightarrow x$ as $n \rightarrow \infty$ (i.e. it converges strongly), then $x_n \rightharpoonup x$.*

This is obvious: $\|x_n - x\| \rightarrow 0$, and any $f \in X^*$ is continuous in the strong topology, so $\|f(x - x_n)\| \rightarrow 0$ too.

Example 1.13.4. The converse is not true: weak convergence does not imply strong convergence. Let $1 < p < \infty$ and consider ℓ^p , with our Schauder basis $\{e^n\}_{n=1}^\infty$. This sequence does not converge strongly, because

$$\|e^n - e^m\|_{\ell^p} = \left(\sum_{i=1}^\infty |e_i^n - e_i^m|^p \right)^{1/p} = 2^{1/p},$$

which does not go to zero. But $e_n \rightarrow 0$! Take any $f \in (\ell^p)^* = \ell^q$; then, $f(e^n) = \sum f_i e_i^n = f_n$, and since $(\sum |f_n|^q)^{1/q}$ is finite, then $|f_n| \rightarrow 0$.

This is clearly a very different kind of convergence than strong convergence, but we have more things converging, which is nice if you like convergence. It’s in some sense a sampling notion.

Proposition 1.13.5. *Let X be an NLS, $\{x_n\} \subset X$ and $\{f_n\} \subset X^*$. Then:*

- *If $\{x_n\}$ converges weakly, then the limit is unique and $\|x_n\|$ is bounded.*
- *If $\{f_n\}$ converges weak-*, then the limit is unique. If in addition X is Banach, then $\|f_n\|$ is bounded.*

PROOF. Suppose $x_n \rightharpoonup x$ and $x_n \rightharpoonup y$. Therefore for any $f \in X^*$, $f(x_n) \rightarrow f(x)$ and $f(x_n) \rightarrow f(y)$. But since limits are unique in \mathbb{F} , then $f(x) = f(y)$, so since this is true for all $f \in X^*$, then $x = y$.

Fix an $f \in X^*$, so that $\{f(x_n)\}_{n=1}^\infty$ converges in \mathbb{F} , and therefore is bounded. Thus, $|E_x(f)| = |f(x_n)| \leq C_f$ for all n (i.e. some constant C that depends only on f). In particular, $\{E_{x_n}\}_{n=1}^\infty \subseteq X^{**}$ is pointwise bounded, so by the uniform boundedness principle, $\|x_n\| = \|E_{x_n}\| \leq C$.

The second part is left as an exercise; but since uniform boundedness requires the space to be Banach, we’ll have to assume that of X . \square

Proposition 1.13.6. *Let X be an NLS and $x_n \rightharpoonup x$. Then, $\|x\| \leq \liminf_{n \rightarrow \infty} \|x_n\|$.*

This one was left as an exercise.

Though most of the time one only cares about convergence, the topologies are worth knowing about. Recall that a topology is a collection of open sets, so one topology being “smaller” than another means that it’s a subset of the other collection of open sets.

Definition. The *weak topology* on X is the smallest topology such that each $f \in X^*$ is continuous. The *weak-* topology* on X^* is the smallest topology such that each E_x (the image of x in the canonical map $X \rightarrow X^{**}$) is continuous.

Since we know these are continuous in the strong topology, these are smaller topologies; there are fewer open sets, and convergence is nicer.

What this means is that if $\mathcal{O} \subset \mathbb{F}$ is open and $f \in X^*$, then $f^{-1}(\mathcal{O})$ is open in X . Since these f are linear, then we really only need to talk about open sets containing 0, and therefore contain an open ball $B_\varepsilon(0)$, which will give us all the open sets in the weak topology (by translation).

We know that if $f_i \in X^*$, then $f_i^{-1}(B_{\varepsilon_i}(0))$ is open, and we can take arbitrary unions and finite intersections. Specifically, our *basic open sets* are of the form

$$\bigcap_{i=1}^n f_i^{-1}(B_{\varepsilon_i}(0)).$$

All open sets in the weak topology are translations of unions of these basic opens. These basic opens are given by sets of the form

$$U = \{x \in X : |f_i(x)| < \varepsilon_i, i = 1, \dots, n\}.$$

For the weak-* topology, the basic opens are

$$\begin{aligned} V &= \{f \in X^* : |f(x_i)| < \varepsilon_i, i = 1, \dots, n\} \\ &= \bigcap_{i=1}^n E_{x_i}^{-1}(B_{\varepsilon_i}(0)). \end{aligned}$$

The next thing we need to do is show that these topologies imply the notions of convergence we defined at the start of lecture.

Lecture 14.

The Banach-Alaoglu Theorem: 9/28/15

“Just take these epsilons; don’t worry about it..”

Recall that we’ve defined weak convergence $x_n \rightarrow x$ if $f(x_n) \rightarrow f(x)$ for all $f \in X^*$, and an even weaker notion called weak-* convergence where $f_n \xrightarrow{w^*} f$ if $f_n(x) \rightarrow f(x)$ for all $x \in X$. Then, we defined the weak and weak-* topologies: the smallest topologies on X (resp. X^*) such that every $f \in X^*$ (resp. $E_x \in X^{**}$) is continuous.

We saw that the basic open sets in the weak topology are finite intersections of sets of the form $f_i^{-1}(B_{\varepsilon_i}(0))$ for $f_i \in X^*$; in other words, a set of points where $f_i(x) < \varepsilon_i$ for some finite number of $f_i \in X^*$ and $\varepsilon_i > 0$. General open sets are unions of translations of these sets. We also had a similar notion for the weak* topology.

Proposition 1.14.1. *Let X be an NLS and $\{x_n\}_{n=1}^\infty \subseteq X$. Then, $x_n \rightarrow x$ in the weak topology iff $x_n \rightarrow x$, and if $\{f_n\}_{n=1}^\infty \subseteq X^*$, then $f_n \rightarrow f$ in the weak* topology iff $f_n \xrightarrow{w^*} f$.*

PROOF. For the first part and in the forward direction, suppose $f \in X^*$, so that f is continuous in the weak topology. Thus, $\lim_{n \rightarrow \infty} f(x_n) = f(x)$, so $x_n \rightarrow x$.

Conversely, suppose $x_n \rightarrow x$, and let U be a basic open set about x , so that

$$U = x + \{y \in X : |f_i(y)| < \varepsilon_i, i = 1, \dots, n\}$$

for some $f_i \in X^*$ and $\varepsilon_i > 0$. Since $f(x_n) \rightarrow f(x)$ for all f , then there’s an $N > 0$ such that $|f_i(x_n - x)| = |f_i(x_n) - f_i(x)| < \varepsilon_i$ for all $i = 1, 2, \dots, n$, and therefore $x_n = (x_n - x) + x \in U$.

The proof of the second part is similar. □

Remark.

- The Hahn-Banach theorem implies that the weak topology is Hausdorff, since it implies there’s a linear functional that separates any two points.
- If τ_w denotes the (collection of open sets of the) weak topology, and τ denotes the strong topology, then $\tau_w \subset \tau$ (i.e. it actually is weaker): since each f_i is continuous in the strong topology, then inverse images of open sets under it remain open. A similar result holds for the weak-* topology.

Theorem 1.14.2 (Banach-Alaoglu¹²). *Let X be an NLS and $B_1^* = \{f \in X^* : \|f\|_{X^*} \leq 1\}$ be the closed unit ball in X^* . Then, B_1^* is compact in the weak-* topology.*

PROOF. Let $B_x = \{\lambda \in \mathbb{F} : |\lambda| \leq |x|\}$. Then, each B_x is a closed and bounded subset of \mathbb{F} , and therefore compact, and let $C = \prod_{x \in X} B_x$. By Tychonoff’s theorem, which is an amazing theorem, any product of compact sets is still compact, so C is compact, even though it’s (in some sense) huge!¹³ A function $g : X \rightarrow \mathbb{F}$ such that $|g(x)| \leq \|x\|$, whether it’s linear or not, can be viewed as an element of C by sending $g \mapsto (g(x))_{x \in X}$. Then, the coordinate map $\pi_x : C \rightarrow B_x$ sends $g \mapsto g(x)$, so it’s just the evaluation map, and we know by the definition of the product topology that this map is continuous. In particular, this defines a continuous inclusion of $B_1^* \subseteq C$, where the former has the weak-* topology. Since C is compact, then B_1^* is compact if it’s closed in C .

We probably won’t get around to showing it, but the weak and weak-* topologies aren’t metrizable, so the nice proof techniques from metric spaces don’t work, and we’ll have to use more basic topological methods. Specifically, we’ll have to show that all accumulation points of B_1^* are in B_1^* .

¹²Pronounced approximately “ala-glue.”

¹³Note that this compactness is in the product topology, not the box topology. If $\{X_\alpha\}_{\alpha \in I}$ is a set of topological spaces and $X = \prod_{\alpha \in I} X_\alpha$, then for each α there’s a coordinate map $\pi_\alpha : X \rightarrow X_\alpha$ sending each element to its coordinate in that index. Then, the product topology is defined to be the smallest topology in which each of these coordinate maps is continuous.

Let g be an accumulation point of B_1^* , and fix $x, y \in X$ and a $\lambda \in \mathbb{F}$; then, let

$$U = g + \{h \in C : |h(x_i)| \leq \varepsilon, i = 1, \dots, m\}.$$

Specifically, we'll take $m = 4$ and the points $x_i = x, y, x + y$, and λx . If $f = g + h \in B_1^*$, we'll choose the epsilons $|h(x)| \leq \varepsilon/3 \max(1, |\lambda|)$, $|h(y)| \leq \varepsilon/3$, $|h(x + y)| \leq \varepsilon/3$, and $|h(\lambda x)| \leq 2\varepsilon/3$, for some arbitrary $\varepsilon > 0$. Then,

$$\begin{aligned} |g(x + y) - g(x) - g(y)| &= |h(x + y) - h(x) - h(y)| \\ &\leq |h(x + y)| + |h(x)| + |h(y)| \leq \varepsilon. \\ |g(\lambda x) - \lambda g(x)| &= h(\lambda x) - \lambda h(x) \leq \varepsilon. \end{aligned}$$

This is why we chose the strange epsilons; as $\varepsilon \rightarrow 0$, this forces g to be linear! Moreover,

$$|g(x)| = |f(x) - h(x)| \leq \|f\| + |h(x)| \leq 1 + \frac{\varepsilon}{3},$$

so in particular $\|g\| \leq 1$, so $g \in B_1^*$, and so B_1^* has to be compact. \square

Now, we can begin to unpack the applications of this theorem. Even though these topologies aren't metrizable, we can get some nice topological results for suitably nice spaces.

Theorem 1.14.3. *Let X be a separable Banach space and $K \subseteq X^*$ be weak-* compact. Then, K with the weak-* topology is metrizable.*

In particular, on this set, sequential compactness is equivalent to compactness, which can be useful!

PROOF Let $D = \{x_n\}_{n=1}^\infty \subseteq X$ be a countable, dense subset of X . Then, the evaluation maps $E_n = E_{x_n}$ are weak-* continuous; since D is dense, then these E_n separate points. That is, if $E_n(x^*) = E_n(y^*)$ for all n (i.e. $x^*(x_n) = y^*(x_n)$ for all n), then $x^* = y^*$.

Let $c_n = \sup_{x^* \in K} |E_n(x^*)|$, which is a continuous, finite function on a compact set and therefore has a maximum, and let

$$f_n = \begin{cases} \frac{E_n}{c_n}, & c_n \neq 0 \\ 0, & c_n = 0. \end{cases}$$

Now that we've scaled suitably, we may define our metric: let

$$d(x^*, y^*) = \sum_{n=1}^\infty 2^{-n} |f_n(x^*) - f_n(y^*)|.$$

This is a metric: it's clearly symmetric, and it's 0 iff each term is, which is true iff $x^* = y^*$, as observed above. Moreover, since the triangle inequality holds termwise, then it holds here.

However, we still need to prove that the topology τ_d induced by this metric agrees with the weak-* topology τ on K . First, we'll show that $\tau_d \subseteq \tau$: for $N \geq 1$, let

$$d_N(x^*, y^*) = \sum_{n=1}^N 2^{-n} |f_n(x^*) - f_n(y^*)|,$$

and consider $d_N(\cdot, y^* : X^* \rightarrow [0, \infty)$ sending $x^* \mapsto d_N(x^*, y^*)$. This is weak-* continuous, and $d_N(\cdot, y^*) \rightarrow d(\cdot, y^*)$ converges uniformly, so the limit $d(\cdot, y^*)$ is continuous. Therefore $B_r(y^*) = \{x^* \in K : d(x^*, y^*) < r\}$ is the inverse image of $(-\infty, r)$ (which is open in \mathbb{R}) under $d(\cdot, y^*)$, which is continuous, so it's open, and thus $\tau_d \subseteq \tau$ (since these balls generate all opens).

Conversely, let $A \in \tau$, so that $A^c \subset K$ is τ -closed, and therefore τ -compact. Since we already showed that $\tau_d \subseteq \tau$, then A^c must be τ_d -compact (as there are fewer open sets), and in particular τ_d -closed. Thus, A is open in τ_d . \square

The Generalized Heine-Borel Theorems: 9/30/15

Recall that we proved the Banach-Alaoglu theorem, that the unit ball in X^* is weak-* compact. This isn't super useful, but we found that if X is a separable Banach space, then B_1^* is a metric space. Relatedly, since any weak-* compact subspace of X^* is metrizable, then if $\{f_n\} \subset X^*$ and $\|f\|_n \leq C$, then there exists a subsequence f_{n_k} which weak-* converges to an $f \in X^*$.

In other words, we know the following.

Theorem 1.15.1 (Generalized Heine-Borel I). *Let X be a separable Banach space and $K \subseteq X^*$. Then, the following are equivalent.*

- (1) K is weak-* compact.
- (2) K is weak-* closed and bounded.
- (3) K is weak-* sequentially compact.¹⁴

We've already seen that (1) \implies (3), and the discussion above is (2) \implies (1) for the unit ball (which generalizes by scaling to any ball, and therefore to any bounded set, since it can be contained in a large ball). So there's only one step left in the proof.

PROOF OF THEOREM 1.15.1, (3) \implies (2). Let K be weak-* sequentially compact. Then, K must be bounded: if not, then there's a sequence $\{f_n\} \subseteq K$ such that $\|f_n\|_{X^*} \geq n$, but then there can be no convergent subsequence, since all weak-* convergent sequences are bounded.

We also want to show that K is weak-* closed, but we don't have a metric topology yet, so we must be careful. Let $\{x_n\}_{n=1}^\infty \subseteq X$ be a countable dense subset and f be an accumulation point of K , so that there's a sequence $f_n \rightarrow f$. In particular, $f_n(x_j) \rightarrow f(x_j)$. Let $U_n = f + \{g \in X^* : |g(x_j)| \leq 1/n, j = 1, 2, \dots, n\}$, which is weak-* open neighborhood of f ; thus, there exists an $f_m \in U_n \cap K$, so that $f = f_m + g$.

If $x \in X$, then there's a subsequence $x_{n_j} \rightarrow x$; we want to know whether $f_n(x) \xrightarrow{?} f(x)$.

$$|f_n(x) - f(x)| \leq |f_n(x) - f_m(x_j)| + |f_m(x_j) - f(x_j)| + |f(x_j) - f(x)|.$$

Each of the three terms on the right goes to zero, but differently (in m or in j). The first one is the only problem, but since it's bounded by $\|f_m\| \|x - x_m\|$ and $\|f_m\| \leq C$, then we may take a large j to make this small, and then the whole expression goes to 0 as $m \rightarrow \infty$, so $f_m \xrightarrow{w^*} f$, and therefore, since K is weak-* sequentially compact, then $f \in K$. \square

We want to take statements about X^* and turn them into statements about x . If X is reflexive (i.e. $X = X^{**}$), then everything we've talked about applies.

Theorem 1.15.2 (Generalized Heine-Borel II). *Let X be a separable, reflexive Banach space and $K \subseteq X$. Then, the following are equivalent.*

- (1) K is compact in the weak topology.
- (2) K is closed in the weak topology and bounded.
- (3) K is sequentially compact in the weak topology.

There's a harder theorem which offers a sort of converse (the Heine-Borel result implies reflexivity), and therefore we can remove the separability hypothesis. However, in applications, one's Banach are almost always separable, and therefore it's not really a huge deal.

Example 1.15.3. Let $1 < p < \infty$ and $\Omega \subseteq \mathbb{R}^d$ be measurable. Then, if $\{f_n\}_{n=1}^\infty \subseteq L^p(\Omega)$ is a sequence such that $\|f_n\|_{L^p} \leq C$, then we know there's a subsequence $f_{n_j} \rightharpoonup f$; that is, for all $g \in L^{p_*}(\Omega)$ (where p_* is the conjugate exponent),

$$\int_{\Omega} f_{n_j}(x)g(x)dx \longrightarrow \int_{\Omega} f(x)g(x)dx.$$

This is used all the time in analysis, and we'll use it later.

We're not done with compactness yet; this next theorem provides a nice connection between weak and strong convergence.

¹⁴Sequential compactness of a set K means that for every sequence $\{x_n\} \subseteq K$, there's a convergent subsequence x_{n_j} .

Theorem 1.15.4 (Banach-Saks). *Let X be an NLS and $x_n \rightarrow x$. Then, for all n , there exist $\alpha_j^{(n)}$ for $j = 1, \dots, n$ with $\alpha_j^{(n)} \geq 0$ and $\sum_{j=1}^n \alpha_j^{(n)} = 1$ such that*

$$y_n = \sum_{j=1}^n \alpha_j^{(n)} x_j \rightarrow x.$$

In other words, there is a convex combination of a weakly convergent sequence that strongly converges. Think about this in the case $e^n \rightarrow 0$.

PROOF. We'll start by considering all such convex combinations. To wit, let

$$M = \left\{ \sum_{j=1}^n \alpha_j^{(n)} x_j : n \geq 1, \alpha_j^{(n)} \geq 0, \sum_{j=1}^n \alpha_j^{(n)} = 1 \right\}.$$

We want to show that $x \in \overline{M}$ (here denoting the strong closure). Notice that M is the convex hull of $\{x_j\}$, and therefore it (and its strong closure) are convex sets.

Let's assume $x \notin \overline{M}$ and use the separating hyperplane theorem (Lemma 1.10.1). We know that \overline{M} is a closed, convex set and $\{x\}$, being a single point, is compact, so there's an $f \in X^*$ and a $\gamma \in \mathbb{R}$ such that $\operatorname{Re} f(x_n) \geq \gamma$, but $\operatorname{Re} f(x) < \gamma$.

Taking the \liminf , we see that $f(x_n) \not\rightarrow f(x)$, and therefore $x_n \not\rightarrow x$, which is a contradiction. \square

Perhaps more interesting is an immediate corollary.

Corollary 1.15.5. *Let X be an NLS and $S \subseteq X$ be convex. Then, the weak and strong closures of S agree.*

PROOF. Let \overline{S}^w denote the weak closure of S , and \overline{S} denote the strong closure. Theorem 1.15.4 tells us that $\overline{S}^w \subseteq \overline{S}$: if $x_n \rightarrow x$, there's a sequence $y_n \rightarrow x$ with $y_n \in S$. Moreover, we already know that $\overline{S} \subseteq \overline{S}^w$, since we already knew that if $x_n \rightarrow x$, then $x_n \rightarrow x$. \square

Dual of an Operator. The dual of an operator goes by many names: dual may be the most common, but it is also known as the *transpose*, *conjugate transpose*, *adjoint*, and so forth.

Definition. Let X and Y be NLSes and $T \in B(X, Y)$. Then, define the *dual* to T to be the map $T^* \in B(Y^*, X^*)$ by $(T^*g)(x) = g(Tx)$.

This makes sense because if $g \in Y^*$ and $x \in X$, then $Tx \in Y$, so $g(Tx) \in \mathbb{F}$. Thus, $T^*g : X \rightarrow \mathbb{F}$. We can write $T^*g = g \circ T$, which writes it as the composition of two continuous (so that T^*g is continuous) and linear (so that T^*g is linear) functions. Thus, $T^*g \in X^*$.

We also have to check that $T^* : Y^* \rightarrow X^*$ is bounded — we've got a perfectly good definition already, but the boundedness of T^* was part of the definition. Suppose $g \in Y^*$ and $x \in X$; then,

$$\begin{aligned} |T^*g(x)| &= |g(Tx)| \leq \|g\|_{Y^*} \|Tx\|_Y \\ &\leq (\|g\|_{Y^*} \|T\|_{B(X,Y)}) \|x\|_X. \end{aligned}$$

Thus, T^* is bounded, and moreover $\|T^*g\| \leq \|g\|_{Y^*} \|T\|_{B(X,Y)}$. Finally, why is T^* linear? By definition,

$$\begin{aligned} T^*(g+h)(x) &= (g+h)(Tx) = g(Tx) + h(Tx) \\ &= (T^*g)(x) + (T^*h)(x) = (T^*g + T^*h)(x) \\ T^*(\lambda g)(x) &= (\lambda g)(Tx) = \lambda g(Tx) = \lambda(T^*g)(x). \end{aligned}$$

We already know that $\|T^*\|_{B(Y^*, X^*)} \leq \|T\|_{B(X,Y)}$; the two are actually equal, and we'll prove this next lecture.

Lecture 16.

The Dual to an Operator: 10/2/15

Recall that last time, we defined the dual to an operator $T \in B(X, Y)$: the dual is $T^* : Y^* \rightarrow X^*$ given by $T^*(g) = g \circ T : X \rightarrow \mathbb{F}$. We showed this is linear and bounded, and in fact $\|T^*g\|_{X^*} \leq \|T\|_{B(X,Y)} \|g\|_{Y^*}$, so taking the supremum, $\|T^*\| \leq \|T\|_{B(X,Y)}$.

Claim. In fact, more is true: $\|T^*\| = \|T\|$.

PROOF By the Hahn-Banach theorem,

$$\begin{aligned}
 \|T\|_{B(X,Y)} &= \sup_{\substack{x \in X \\ x \neq 0}} \frac{\|Tx\|_Y}{\|x\|_X} = \sup_{\substack{x \in X \\ x \neq 0}} \sup_{\substack{g \in Y^* \\ g \neq 0}} \frac{|g(Tx)|}{\|x\| \|g\|} \\
 &\leq \sup_x \sup_g \frac{\|T^*g\| \|x\|}{\|x\| \|g\|} \\
 &\leq \sup_g \frac{\|T^*\| \|g\|}{\|g\|} = \|T^*\|. \quad \square
 \end{aligned}$$

We have another operator around, $*$: $B(X, Y) \rightarrow B(Y^*, X^*)$ sending $T \mapsto T^*$. One can show that it's linear and bounded, though this is left as an exercise.

Proposition 1.16.1.

- (1) $*$ is an isometry: $\|T\| = \|T^*\|$.
- (2) $*$ is linear: $(\lambda T + \mu S)^* = \lambda T^* + \mu S^*$.
- (3) $*$ is contravariant: if $S : X \rightarrow Y$ and $R : Y \rightarrow Z$, then $(R \circ S)^* = S^* \circ R^*$ as maps $Z^* \rightarrow Y^* \rightarrow X^*$.
- (4) $(\text{id}_X)^* = \text{id}_{X^*}$.

This is also left as an exercise, though maybe it reminds you of something: transposes of matrices.¹⁵ For example, if $T : \mathbb{R}^m \rightarrow \mathbb{R}^n$, then T is represented by an $(n \times m)$ -matrix A ($T(x) = Ax$). Then, $T^* = A^T$; showing this is a useful exercise.

Example 1.16.2. Let $1 < p < \infty$ and $f \in L^p(0, 1)$. For any $x \in (0, 1)$, define an *integral operator* $T : L^p(0, 1) \rightarrow L^p(0, 1)$ by

$$Tf(x) = \int_0^1 K(x, y)f(y) dy,$$

where $K \in L^\infty((0, 1) \times (0, 1))$. We know that the dual space is $(L^p)^* = L^q$, where $1/p + 1/q = 1$, in the sense that $g \in L^q$ iff the map $\Lambda_g \in (L^p)^*$, where

$$\Lambda_g(f) = \int_0^1 g(x)f(x) dx.$$

Let's find out what T^* is: this should be a map $(L^p)^* \rightarrow (L^p)^*$. We know

$$\begin{aligned}
 (T^* \Lambda_g)(f) &= \Lambda_g(Tf) = \int_0^1 g(x)Tf(x) dx \\
 &= \int_0^1 g(x) dx \int_0^1 K(x, y)f(y) dy dx.
 \end{aligned}$$

Using the Fubini theorem, we may change the order of integration, because we want to recast this as an operator in terms of f .

$$\begin{aligned}
 &= \int_0^1 f(y) \int_0^1 K(x, y)g(x) dx dy \\
 &= \int_0^1 f(x) \underbrace{\int_0^1 K(y, x)g(y) dy}_{T^* \Lambda_g} dx.
 \end{aligned}$$

That is, $T^* \Lambda_g = \Lambda_{\int_0^1 K(y, x)g(y) dy}$. Once again, $K(x, y)$ being sent to $K(y, x)$ is a sort of transposition.

Whenever we can do something once, we like to do it twice.

Lemma 1.16.3. Let X and Y be NLSes and $T \in B(X, Y)$. Then, $T^{**} \in B(X^{**}, Y^{**})$ is a bounded extension of T , i.e. under the canonical inclusion $X \hookrightarrow X^{**}$, $T^{**}|_X = T$.

¹⁵What it reminds you of definitely depends on who you are; I see a contravariant functor!

If X is reflexive, this means that $T^{**} = T$. For example, in finite-dimensional vector spaces, applying the transpose twice gets you back where you started.

PROOF. Let $x \in X$ and $g \in Y^*$. Then, $T^{**}(E_x)(g) = E_x(T^*g) = (T^*g)(x) = g(Tx) = E_{Tx}(g)$, i.e. $T^{**}E_x = E_{Tx}$. \square

The proof amounts to unwinding definitions.

Back in the world of matrices, if the matrix is square we can sometimes take the inverse; remember that this commutes with the transpose. Let's generalize this.

Lemma 1.16.4. *Let X be a Banach space, Y be an NLS, and $T \in B(X, Y)$. Then, T has a bounded inverse on Y iff T^* has a bounded inverse on all of X^* ; in this case, $(T^*)^{-1} = (T^{-1})^*$.*

PROOF. First, the forward direction: suppose $S = T^{-1} \in B(Y, X)$. Then, $S^*T^* = (TS)^* = (I_Y)^* = I_{Y^*}$ (where I_Y is the identity on Y), and therefore T^* is one-to-one. If we go the other way, $T^*S^* = (ST)^* = (I_X)^* = I_{X^*}$, so T^* is onto. Thus, $(T^*)^{-1}$ exists (as a linear map), and these calculations showed us that it's S^* , so in particular, $(T^*)^{-1}$ is bounded.

In the other direction, we know $(T^*)^{-1}$ exists, and therefore $(T^{**})^{-1}$ exists, and in particular is bijective, and $T^{**}|_X = T$, so in particular T is one-to-one (since T^{**} is), so we need to show that it maps onto Y , because then the open mapping theorem will imply it has a bounded inverse.

We know T^{**} maps onto and is an open map, so it takes open sets to open sets, and therefore closed sets to closed sets. Since X is Banach and therefore closed in X^{**} , then $T^{**}(X) = T(X)$ is closed in Y^{**} , and therefore $T(X)$ is closed in Y . Suppose T isn't onto Y , so that there exists a $y \in Y \setminus T(X)$, and the Hahn-Banach theorem allows us to create a $g \in Y^*$ such that $g|_{T(X)} = 0$ but $g(y) = \|y\| \neq 0$. But then we see that if $x \in X$, $(T^*g)(x) = g(Tx) = 0$, but since T^* is invertible, this means $g = 0$, and therefore $\|y\| = 0$, so $y = 0$. But $0 = T(0)$, so this means T is onto. \square

So ends Chapter 2 of the book. Chapter 3 is easier: many introductions to functional analysis start with Hilbert spaces, which are easier, and then ramp it up to where we were. And some abstract introductions start with a very abstract notion of a locally convex vector space!

A Hilbert space is a Banach space, but with more structure. They solve the problem that, though we have a nice notion of size, we don't have a notion of angle, like in finite-dimensional vector spaces. Hilbert spaces have a solution to this.

Definition. An *inner product* on a vector space H , denoted (\cdot, \cdot) , $(\cdot, \cdot)_H$, $\langle \cdot, \cdot \rangle$, or $\langle \cdot, \cdot \rangle_H$ is a map $H \times H \rightarrow \mathbb{F}$ such that:

- (1) (\cdot, \cdot) is linear in its first argument.
- (2) (\cdot, \cdot) is *conjugate symmetric*, i.e. $(x, y) = \overline{(y, x)}$: reversing the arguments produces the complex conjugate. If \mathbb{F} is real, then this means that (\cdot, \cdot) is symmetric.
- (3) For any $x \in H$, $(x, x) \geq 0$, and $(x, x) = 0$ iff $x = 0$.

H along with this inner product is called an *inner product space*, which we'll abbreviate IPS.

If you combine properties (1) and (2), one sees that (\cdot, \cdot) is *conjugate linear* in its second argument:

$$\begin{aligned} (x, \alpha y + \beta z) &= \overline{(\alpha y + \beta z, x)} \\ &= \overline{\alpha(y, x) + \beta(z, x)} \\ &= \overline{\alpha}(y, x) + \overline{\beta}(z, x) \\ &= \overline{\alpha}(x, y) + \overline{\beta}(x, z). \end{aligned}$$

This property, also known as *sesquilinearity*, means that it commutes with addition, but scalar multiplication in the second argument gets replaced with its conjugate.

Example 1.16.5.

- (1) \mathbb{F}^d is an IPS with the *complex dot product*

$$(x, y) = x \cdot \overline{y} = \sum_{i=1}^d x_i \overline{y_i}.$$

It's not hard to show that this satisfies the three defining properties.

- (2) ℓ^p is *not* an inner product space, unless $p = 2$. The inner product on ℓ^2 is

$$(x, y) = x \cdot \bar{y} = \sum_{i=1}^{\infty} x_i \bar{y}_i.$$

By the Hölder inequality, this is bounded by $\|x\|_{\ell^2} \|\bar{y}\|_{\ell^2} = \|x\|_{\ell^2} \|y\|_{\ell^2}$, so the inner product is finite, and has the desired properties.

- (3) Similarly, $L^p(\Omega)$ isn't an inner product space unless $p = 2$; an inner product on $L^2(\Omega)$ is given by

$$(f, g) = \int_{\Omega} f(x) \overline{g(x)} dx.$$

The idea here is that 2 is its own conjugate exponent, because $1/2 + 1/2 = 1$. Hölder's inequality once again guarantees that this is finite.

Inner Product Spaces and Hilbert Spaces

Lecture 17.

Orthogonality: 10/5/15

Recall that last time, we defined an inner product $(\cdot, \cdot) : X \times X \rightarrow \mathbb{F}$ on a vector space X , which is linear in the first argument, satisfies $(x, y) = \overline{(y, x)}$, and is positive definite: $(x, x) \geq 0$, and is equal to 0 iff $x = 0$. Then, X is called an inner product space (IPS).

An inner product buys us a lot of other structure, too.

Definition. The *induced norm* on an IPS $(H, (\cdot, \cdot))$ is the function $\|\cdot\| : H \rightarrow [0, \infty)$ defined by $\|x\| = (x, x)^{1/2}$.

This is, unsurprisingly, a norm: first, $\|x\| \geq 0$ and is 0 iff $x = 0$ by positive definite-ness. Since the inner product is linear, then $\|\lambda x\| = |\lambda| \|x\|$, because

$$\|\lambda x\|^2 = (\lambda x, \lambda x) = \lambda \bar{\lambda} (x, x) = |\lambda|^2 \|x\|^2.$$

Using the norm squared sometimes allows us to save writing some square roots.

For the triangle inequality we need an intermediate result, analogous to Hölder's inequality with $p = 2$.

Lemma 2.1.1 (Cauchy-Schwarz inequality). *If H is an IPS and $x, y \in H$, $|(x, y)| \leq \|x\| \|y\|$.*

PROOF. We're done if $y = 0$, so assume $y \neq 0$ and $\lambda \in \mathbb{F}$. Then,

$$\begin{aligned} 0 &\leq \|x - \lambda y\|^2 = (x - \lambda y, x - \lambda y) \\ &= (x, x - \lambda y) - \lambda (x, x - \lambda y) \\ &= (x, x) - \bar{\lambda} (x, y) - \lambda (y, x) + |\lambda|^2 (y, y) \\ &= \|x\|^2 - 2 \operatorname{Re}(\lambda (y, x)) + |\lambda|^2 \|y\|^2. \end{aligned} \tag{2.1.1}$$

This is a positive, quadratic function in λ , so by taking the derivative with respect to λ , the minimum is $\lambda = (x, y) / \|y\|^2$. Choose this λ , so as to obtain the maximum amount of information. Thus, in this case, (2.1.1) becomes

$$\begin{aligned} 0 &\leq \|x\|^2 - 2 \frac{|(x, y)|^2}{\|y\|^2} + \frac{|(x, y)|^2}{\|y\|^4} \|y\|^2 \\ &= \|x\|^2 - \frac{|(x, y)|^2}{\|y\|^2}. \end{aligned} \quad \square$$

The proof reminds us of Hölder's inequality, too. However, since the inner product is sesquilinear, it's important to keep track of complex conjugates to avoid adding errors.

Now, we can prove the triangle inequality for the induced norm.

$$\begin{aligned} \|x + y\|^2 &= (x + y, x + y) \\ &= \|x\|^2 + \|y\|^2 + 2 \operatorname{Re}(x, y) \\ &\leq \|x\|^2 + \|y\|^2 + 2|(x, y)| \\ &\leq \|x\|^2 + \|y\|^2 + 2\|x\| \|y\| \\ &= (\|x\| + \|y\|)^2. \end{aligned}$$

In other words, *inner product spaces are normed spaces too*.

The Cauchy-Schwarz inequality also suggests to us that we have a well-defined notion of angle.

Definition.

- If H is an inner product space over $\mathbb{F} = \mathbb{R}$, we can define the *angle* θ between two points $x, y \in H$ as the solution to

$$\cos \theta = \frac{(x, y)}{\|x\| \|y\|} \in [-1, 1].$$

- If H is a complex IPS, we define the *angle* between $x, y \in H$ to be the θ such that

$$\cos \theta = \frac{|(x, y)|}{\|x\| \|y\|} \in \left[0, \frac{\pi}{2}\right].$$

- If $(x, y) = 0$, so that the angle between them is $\pi/2$, then x and y are said to be *orthogonal*, written $x \perp y$.

Though we only have a reduced notion of angle in complex vector spaces, it's OK, because we mostly care about orthogonality.

We also have a nice formula for addition: x, y , and 0 are three vertices of a parallelogram, and $x + y$ is the fourth vertex.

Proposition 2.1.2 (Parallelogram law). $\|x + y\|^2 + \|x - y\|^2 = 2(\|x\|^2 + \|y\|^2)$.

Note that this is not true for all norms; in particular, if you can find one that doesn't satisfy the parallelogram law, then you've shown that not every NLS is an IPS.

We can also talk about topology.

Lemma 2.1.3. $(\cdot, \cdot) : X \times X \rightarrow \mathbb{F}$ is continuous.

Corollary 2.1.4. Suppose $\lambda_n \rightarrow \lambda$ and $\mu_n \rightarrow \mu$ in \mathbb{F} , and that $x_n \rightarrow x$ and $y_n \rightarrow y$ in H . Then, $(\lambda_n x_n, \mu_n y_n) \rightarrow (\lambda x, \mu y)$.

PROOF OF LEMMA 2.1.3. To avoid confusion between the inner product and elements of the product space, we'll use $\langle \cdot, \cdot \rangle$ to denote the inner product on H in this proof.

Since H (and thus also $H \times H$) and \mathbb{F} are metric spaces, continuity is equivalent to sequential continuity, so suppose $(x_n, y_n) \rightarrow (x, y)$ in $H \times H$, i.e. $x_n \rightarrow x$ and $y_n \rightarrow y$ in H . Thus,

$$\begin{aligned} |\langle x_n, y_n \rangle - \langle x, y \rangle| &\leq |\langle x_n, y_n \rangle - \langle x_n, y \rangle| + |\langle x_n, y \rangle - \langle x, y \rangle| \\ &= |\langle x_n, y_n - y \rangle| + |\langle x_n - x, y \rangle| \\ &\leq \|x_n\| \|y_n - y\| + \|x_n - x\| \|y\|. \end{aligned}$$

Since $\{x_n\}$ and $\{y_n\}$ are convergent sequences, then their norms are bounded, so we're done. \square

Then, Corollary 2.1.4 follows directly: we know that scalar multiplication is continuous in any vector space. Lemma 2.1.3 shouldn't come as a surprise: we've defined a continuous norm, after all.

Best Approximation and Orthogonal Projection.

Definition. Just as a Banach space is a complete normed space, a *Hilbert space* is a complete IPS.

Theorem 2.1.5 (Best approximation). Let $(H, (\cdot, \cdot))$ be an IPS and $M \subseteq H$ be a nonempty, convex, and complete¹ subset. If $x \in H$, then there exists a unique $y \in M$ such that

$$\text{dist}(x, M) = \inf_{z \in M} \|x - z\| = \|x - y\|.$$

In other words, y is the best approximation (minimizing distance) to x that's in M .

PROOF. Let $\delta = \inf_{z \in M} \|x - z\|$. If $\delta = 0$, then $y = x$, and moreover $x \in M$, because M is complete and therefore closed. This tells us that the best approximation to a point in the space is itself, which is perhaps unsurprising.

¹If H is a Hilbert space, then completeness is the same as being closed.

If $\delta > 0$, then $x \neq M$, so there exists a sequence $\{y_n\}_{n=1}^{\infty}$ such that $\delta_n = \|x - y_n\| \rightarrow \delta$. It's easy to see that $\{y_n\}$ is Cauchy: by Lemma 2.1.2,

$$\begin{aligned}\|y_n - y_m\|^2 &= \|(y_n - x) + (x - y_m)\| \\ &= 2(\|y_n - x\|^2 + \|x - y_m\|^2) - \|y_n + y_m - 2x\| \\ &= 2(\delta_n^2 + \delta_m^2) - 4\left\|\frac{y_n + y_m}{2} - x\right\|^2.\end{aligned}$$

Since M is convex, $(y_n + y_m)/2 \in M$, and therefore

$$\leq 2(\delta_n^2 + \delta_m^2) - 4\delta^2 \rightarrow 0.$$

Since M is complete and $\{y_n\}$ is Cauchy, then $y_n \rightarrow y$, which therefore attains the infimum and is our best approximation. Now, we have to prove uniqueness: if $\|x - z\| = \delta$, then

$$\|y - z\|^2 = 4\delta^2 - 4\left\|\frac{y + z}{2} - x\right\|^2 \leq 4\delta^2 - 4\delta^2 = 0. \quad \square$$

Corollary 2.1.6. If $M \subset H$ is a complete linear subspace, $x \in H$, and $y \in M$ is its best approximation, then $x - y \perp M$.²

PROOF. Let $m \in M$ be nonzero. Then, for any $\lambda \in \mathbb{F}$,

$$\|x - y\|^2 \leq \|x - y + \lambda m\|^2 = \|x - y\|^2 + |\lambda|^2 \|m\|^2 + 2\operatorname{Re} \bar{\lambda}(x - y, m).$$

Taking $\lambda = -(x - y, m)/\|m\|^2$, we conclude that

$$\begin{aligned}0 &\leq |\lambda|^2 \|m\|^2 - \frac{2|(x - y, m)|}{\|m\|^4} \|m\|^2 \\ &= |\lambda|^2 \|m\|^2 - 2|\lambda|^2 \|m\|^2 = -\lambda^2 \|m\|^2,\end{aligned}$$

but this only makes sense when $\lambda = 0$, which means that $(x - y, m) = 0$. \(\square\)

This is nothing more than projection from a vector space down to a subspace (though we do require completeness in this case). It'll be useful to consider all of the points which project down.

Definition. The *orthogonal complement* of an $M \subseteq H$ is the set $M^\perp = \{x \in H : x \perp M\}$ (i.e. the set of x such that $(x, m) = 0$ for all $m \in M$).

For example, \mathbb{R}^2 is a Hilbert space with its usual inner product; if M is the x -axis, then M^\perp is the y -axis. As in that case, we'd more generally want to write $H = M + M^\perp$, but we'll do that next time.

Lecture 18.

Projections: 10/7/15

Today's the exam, from 7 to 9 pm in UTC 1.104; it will cover §2.1 – 2.7 from the textbook.

We were talking about best approximation: if H is an IPS, $M \subseteq H$ is a nonempty, convex, and complete subset, and $x \in H$, then there exists a best approximation, i.e. a unique $y \in M$ such that $\operatorname{dist}(x, M) = \|x - y\|$. As a corollary, if M is a complete linear subspace (which is automatically nonempty and convex), we know $x - y \perp M$, so if $\langle \cdot, \cdot \rangle$ denotes the inner product, then $\langle x - y, m \rangle = 0$ for all $m \in M$. This motivated the definition of the orthogonal complement M^\perp , the set of all vectors orthogonal to M .

Proposition 2.2.1. If H is an IPS and $M \subseteq H$, then M^\perp is a closed linear subspace of H ; furthermore, $M \perp M^\perp$ and $M \cap M^\perp = \{0\}$ or is empty.

This will be left as an exercise, because it's not very hard. That M^\perp is closed ultimately follows from the continuity of the inner product.

The best approximation is a projection, a concept you might recall from finite-dimensional linear algebra.

Definition. If X is an NLS and $P : X \rightarrow X$ satisfies $P^2 = P$, then P is called a *projection*.

² $x \perp v$ if for all $v \in V$, $x \perp v$, just as in linear algebra.

For example, on \mathbb{R}^2 , we could project something down to the x -axis; if you do this twice, you're still on the x -axis, and nothing more changes. Also, if P is a projection and $M = \text{Im}(P)$, then $P|_M = \text{id}$: if $m \in M$, then $m = P(x)$ for some x , so $P(m) = P^2(x) = P(x) = m$.

Proposition 2.2.2. *Let X be an NLS and $P : X \rightarrow X$ be a projection mapping onto $M \subseteq X$. If $Q = I - P$, then Q is a projection, and $QP = PQ = 0$.*

PROOF. This is just algebra:

$$Q^2 = (I - P)(I - P) = I^2 - IP - PI + P^2 = I - 2P + P = I - P = Q,$$

since P is a projection, so $P^2 = P$. Then, $QP = (I - P)P = P - P^2 = 0$, and $PQ = P(I - P) = P - P^2 = 0$. \square

We're particularly interested in projections that are also linear operators.

Proposition 2.2.3. *Let P and Q be as in Proposition 2.2.2, where P projects onto M and Q projects onto N . If P is linear, then so is Q , and $X = M \oplus N$.³ If P is bounded and $M \neq \{0\}$, then $\|P\| \geq 1$.*

PROOF. Since $Q = I - P$ is a difference of linear functions, then it's also linear.

To show that $X = M \oplus N$, we must show that $X = M + N$ and $M \cap N = \{0\}$. For the first claim, any $x \in X$ can be written as $x = x - Px + Px = (I - P)x + Px = Qx + Px$, and $Qx \in N$ and $Px \in M$; for the second, if $x \in M \cap N$, then $Px = x$ and $Qx = x$, so $PQx = x$, but $PQ = 0$, so $x = 0$.

Finally, for any $m \in M \setminus \{0\}$,

$$\|P\| = \sup_{x \in X \setminus \{0\}} \frac{\|Px\|}{\|x\|} \geq \frac{\|Pm\|}{\|m\|} = \frac{\|m\|}{\|m\|} = 1. \quad \square$$

The last part of this proposition is interesting: there are linear projections where the norm could increase; this is a little counterintuitive.

In an inner product space, we also have a notion of orthogonal projection.

Definition. Let H be an IPS and $M \subseteq H$ be a complete linear subspace. Then, define $P = P_M : X \rightarrow M$ as sending X to its best approximation in M , and define $P^\perp = P_M^\perp$ to be $I - P$.

By Theorem 2.1.5, Px (and therefore also $P^\perp x$) is uniquely defined for every $x \in X$.

Lemma 2.2.4. *P is a projection onto M and P^\perp is a projection onto M^\perp .*

PROOF. We proved that the best approximation in M of an $x \in M$ is just x again, so P^2x is $P(Px)$, but since $Px \in M$, this is just Px , so $P^2 = P$. By Proposition 2.2.2, P^\perp is a projection as well, and by Corollary 2.1.6, for any $x \in X$, $(I - P)x \in M^\perp$. Then, if $x \in M^\perp$, then its best approximation $Px \perp M$, but $Px \in M$ as well, so $Px = 0$, and therefore $x = P^\perp x$. Thus, P^\perp is onto M^\perp . \square

By Proposition 2.2.2, this means $PP^\perp = P^\perp P = 0$ and $P^2 = (P^\perp)^\perp = 0$; Lemma 2.2.4 also implies that $P_M^\perp = P_{M^\perp}$. Additionally, $x \in M$ iff $x = Px$ iff $P^\perp x = 0$, and thus also $x \in M^\perp$ iff $x = P^\perp x$ iff $Px = 0$.

Theorem 2.2.5. *P and P^\perp are bounded linear operators.*

PROOF. Let $x, y \in X$ and $\alpha, \beta \in \mathbb{F}$; then,

$$\begin{aligned} \alpha x + \beta y &= P(\alpha x + \beta y) + P^\perp(\alpha x + \beta y) \\ &= \alpha(Px + P^\perp x) + \beta(Py + P^\perp y), \end{aligned}$$

and therefore

$$\alpha Px + \beta Py - P(\alpha x + \beta y) = P^\perp(\alpha x + \beta y) - \alpha P^\perp x - \beta P^\perp y, \quad (2.2.1)$$

so (2.2.1) lies in both M (on the left) and M^\perp (on the right), so it's equal to 0. Thus, we can conclude that $P(\alpha x + \beta y) = \alpha Px + \beta Py$, and similarly for P^\perp in place of P , so they're both linear.

It suffices to show that P is bounded, because then it follows that P^\perp is too. Since $P + P^\perp = I$, then

$$\begin{aligned} \|x\|^2 &= \|Px + P^\perp x\|^2 = \langle Px + P^\perp x, Px + P^\perp x \rangle \\ &= \|Px\|^2 + \underbrace{\langle Px, P^\perp x \rangle}_0 + \underbrace{\langle P^\perp x, Px \rangle}_0 + \|P^\perp x\|^2. \end{aligned} \quad (2.2.2)$$

That is, $\|Px\|^2 = \|x\|^2 - \|P^\perp x\|^2 \leq \|x\|^2$, meaning P is bounded and $\|P\| \leq 1$. \square

³If two subspaces $M, N \subseteq X$ are such that $X = M + N$ and $M \cap N = \{0\}$, then one says that $X = M \oplus N$, the *direct sum* of M and N .

(2.2.2) has a familiar-looking corollary.

Corollary 2.2.6 (Pythagorean theorem). *For any $x \in X$, $\|x\|^2 = \|Px\|^2 + \|P^\perp x\|^2$.*

It also leads us to the the next result.

Corollary 2.2.7. *If $M \neq 0$, then $\|P\| = 1$, and if $M \neq X$, then $\|P^\perp\| = 1$.*

PROOF. The proof of Theorem 2.2.5 shows $\|P\| \leq 1$, and Proposition 2.2.3 says that if $M \neq 0$, then $\|P\| \geq 1$, so in this case $\|P\| = 1$. Since $M^\perp \neq 0$ iff $M \neq X$, then in this case $\|P^\perp\| = 1$ too. \square

These results (particularly that $P_M^\perp = P_{M^\perp}$) are why the best approximation is often called the *orthogonal projection*, and we will adopt this term. The following result, which is a useful characterization of Px , is sometimes taken as its definition.

Proposition 2.2.8. *Let X , M , and P be as above and $x, y \in X$; then, $y = Px$ iff $y \in M$ and $\langle x - y, m \rangle = 0$ for all $m \in M$.*

PROOF. The forward direction is already established, so suppose $y \in M$ and $\langle x - y, m \rangle = 0$ for all $m \in M$. Then, $\|x - (y + m)\|^2 = \|x - y\|^2 + \|m\|^2$ is minimal over $m \in M$ for the best approximation Px , but is also minimal if $m = 0$ (i.e. for y), so $Px = y$. \square

Dual Spaces. Let H be a Hilbert space; then, we'll define some elements of the dual space. For example, if $y \in H$, then let $L_y : H \rightarrow \mathbb{F}$ be given by $L_y(x) = \langle x, y \rangle$. This is linear and bounded: by the Cauchy-Schwarz inequality, $|L_y(x)| = |\langle x, y \rangle| \leq \|x\| \|y\|$, so $\|L_y\|_{H^*} \leq \|y\|$. But since

$$L_y\left(\frac{y}{\|y\|}\right) = \left\langle \frac{y}{\|y\|}, y \right\rangle = \|y\|,$$

then $\|L_y\|_{H^*} = \|y\|$.

It turns out this is everything.

Theorem 2.2.9 (Riesz representation theorem). *If $(H, \langle \cdot, \cdot \rangle)$ is a Hilbert space and $L \in H^*$, then there's a unique y such that $L = L_y$ (i.e. $Lx = \langle x, y \rangle$), and so $\|L\|_{H^*} = \|y\|_H$.*

PROOF. The proof is very clever, and uses a trick.

First, let's use uniqueness: suppose $Lx = \langle x, y_1 \rangle = \langle x, y_2 \rangle$. Then, $\langle x, y_1 - y_2 \rangle = 0$ for every $x \in X$, but this means that $y_1 - y_2 = 0$ (test $x = y_1 - y_2$; we know $\|y_1 - y_2\|^2 = 0$ iff $y_1 = y_2$).

For existence, first note that if $L = 0$, we can choose $y = 0$. Thus, suppose $L \neq 0$, and let $M = \ker(L)$, i.e. $M = \{x \in H : Lx = 0\}$. Since $L \neq 0$, then $M \subsetneq H$. Since $M = L^{-1}(0)$ and $\{0\}$ is closed, then M is closed.

Now for the weird part: choose any $z \in M^\perp$ such that $\|z\| = 1$, and consider $u = (Lx)z - (Lz)x$, so that $Lu = 0$, and thus $u \in M$. In particular, $u \perp z$, so

$$0 = \langle u, z \rangle = \langle (Lx)z - (Lz)x, z \rangle = (Lx)\|z\|^2 - Lz(x, z)$$

and so $Lx = (x, (Lz)z)$, so let $y = (Lz)z$. \square

Lecture 19.

Orthonormal Bases: 10/9/15

“Looks like a thinner class after the exam...”

In a finite-dimensional space, you can produce a basis of orthogonal vectors. It turns out you can do this in Hilbert spaces as well; our next goal is to prove this.

Definition.

- Let H be an IPS and \mathcal{J} be an index set. Then, a set $A = \{x_\alpha\}_{\alpha \in \mathcal{J}} \subset H$ is *orthogonal* if $x_\alpha \neq 0$ for all $\alpha \in \mathcal{J}$ and $x_\alpha \perp x_\beta$ for all $\alpha, \beta \in \mathcal{J}$ such that $\alpha \neq \beta$ (they're pairwise orthogonal). If in addition $\|x_\alpha\| = 1$ for all α , A is said to be *orthonormal* (sometimes abbreviated ON).
- Let X be an NLS (no inner product needed here) and $A \subseteq X$. Then, A is *linearly independent* if every finite subset of A is linearly independent, i.e. if $\{x_i\}_{i=1}^n \subseteq A$ and $c_1x_1 + \dots + c_nx_n = 0$, then $c_1 = \dots = c_n = 0$.

Proposition 2.3.1. *Let H be an inner product space and $A \subseteq H$ be an orthogonal subset. Then, A is linearly independent.*

PROOF. This is essentially the same proof as one does in finite dimensions: let $\{x_i\}_{i=1}^n$ be any finite subset of A and $c_i \in \mathbb{F}$ be such that $\sum_{i=1}^n c_i x_i = 0$. For any j , (\cdot, x_j) is linear, so

$$0 = \sum_{i=1}^n c_i (x_i, x_j) = c_j \|x_j\|^2.$$

Since $x_j \neq 0$, then $c_j = 0$. \(\square\)

In order to talk about bases, we'll need to talk about projections again. Suppose that $\{x_1, \dots, x_n\}$ is a linearly independent subset of H , so that $M = \text{span}\{x_1, \dots, x_n\}$ is a closed subspace of H . For any $x \in X$, $P_M x \in M$, so there exist $c_1, \dots, c_n \in \mathbb{F}$ such that

$$P_M x = \sum_{j=1}^n c_j x_j. \quad (2.3.1)$$

Let's calculate these c_j : since $P_M x - x \perp M$ by Proposition 2.2.8, then $(P_M x, x_i) = (x, x_i)$ for all i . Let $\mathbf{c} = (c_1, \dots, c_n)$, A be the matrix whose entries are $a_{ij} = (x_j, x_i)$, and $\mathbf{b} = (b_1, \dots, b_n)$, where $b_i = (x, x_i)$. Thus, (2.3.1) means that $A\mathbf{c} = \mathbf{b}$, so we can recover the c_j coefficients by $\mathbf{c} = A^{-1}\mathbf{b}$, assuming A is invertible.

This may be harder to compute in general, but orthogonality comes to our assistance: $A = I$, so $\mathbf{c} = \mathbf{b}$. In other words, we've proven (2.3.2a) in the following theorem.

Theorem 2.3.2. *Let H be a Hilbert space and $\{u_1, \dots, u_n\} \subseteq H$ be an orthonormal set. Let $M = \text{span}\{u_1, \dots, u_n\}$ and $x \in H$; then,*

$$P_M x = \sum_{i=1}^n (x, u_i) u_i, \quad (2.3.2a)$$

and

$$\sum_{i=1}^n |(x, u_i)|^2 \leq \|x\|^2. \quad (2.3.2b)$$

PROOF. We've seen the proof of (2.3.2a), so let's look at (2.3.2b). Since P_M is an orthogonal projection, then

$$\|P_M x\|^2 \leq \|x\|^2 = \left(\sum_i (x, u_i) u_i, \sum_j (x, u_j) u_j \right) = \sum_i |(x, u_i)|^2. \quad \square$$

This proof once again looks a lot like what we did in finite-dimensional linear algebra. So let's do something that doesn't.

Definition. Let \mathcal{J} be any index set, possibly uncountable, and choose $\{x_\alpha\}_{\alpha \in \mathcal{J}}$ with $x_\alpha \in [0, \infty)$. Then, define their sum to be

$$\sum_{\alpha \in \mathcal{J}} x_\alpha = \sup_{\substack{J \subseteq \mathcal{J} \\ J \text{ finite}}} \sum_{\beta \in J} x_\beta.$$

For example, if $\mathcal{J} = \mathbb{N}$, then this looks familiar:

$$\sum_{\alpha=0}^{\infty} x_\alpha = \lim_{n \rightarrow \infty} \sum_{\alpha=0}^n x_\alpha.$$

However, adding uncountability doesn't get us very much: it turns out that if $\sum_{\alpha \in \mathcal{J}} x_\alpha$ is finite, then at most countably many x_α are nonzero! This is worth thinking about, but it makes sense: only finitely many can be greater than ε for any $\varepsilon > 0$.

Theorem 2.3.3 (Bessel's inequality). *let H be a Hilbert space and $\{u_\alpha\}_{\alpha \in \mathcal{J}} \subseteq H$ be an orthonormal set. Then, for any $x \in H$,*

$$\sum_{\alpha \in \mathcal{J}} |(x, u_\alpha)|^2 \leq \|x\|^2.$$

PROOF. If $J \subseteq \mathcal{J}$ is a finite set, then Theorem 2.3.2 tells us that

$$\sum_{\beta \in J} |(x, u_\beta)|^2 \leq \|x\|^2,$$

so this is still true when we take the supremum. \(\square\)

Corollary 2.3.4. *At most finitely many (x, u_α) are nonzero.*

This might be a little strange: no matter how large this Hilbert space is, every vector can only project down to finitely many vectors in an orthonormal set.

We're working up to having an orthonormal basis for a Hilbert space, so let's consider some examples. \mathbb{F}^d is a Hilbert space with $\|\cdot\|_{\ell^2}$ induced by the usual inner product, so we can take an indexing set $\mathcal{J} = \{1, \dots, d\}$. ℓ^2 with the usual norm and inner product (this time given by an infinite sum) can take the indexing set $\mathcal{J} = \mathbb{N}$ (starting from 1, not 0), which produces an orthonormal basis. We want to generalize this to possibly uncountable index sets \mathcal{J} , producing larger Hilbert spaces which are in some sense indexed in \mathcal{J} ; then, the inner product should be similar, but with the sum over \mathcal{J} . Corollary 2.3.4 implies that such a large sum still makes sense.

Let's make this formal.

Definition. Let \mathcal{J} be an index set, and let

$$\ell^2(\mathcal{J}) = \left\{ f : \mathcal{J} \rightarrow \mathbb{F} : \sum_{\alpha \in \mathcal{J}} |f(\alpha)|^2 \text{ is finite} \right\}.$$

Often, we'll let f_α denote $f(\alpha)$.

We'll end up proving that *all* Hilbert spaces are isomorphic to some $\ell^2(\mathcal{J})$! All you really need to know about a Hilbert space is how big it is.

Theorem 2.3.5 (Riesz-Fisher). *Let H be a Hilbert space and $\{u_\alpha\}_{\alpha \in \mathcal{J}} \subseteq H$ be orthonormal. Define $F : H \rightarrow \ell^2(\mathcal{J})$ by $F(x) = f_x$, where $f_x(\alpha) = x_\alpha = (x, u_\alpha)$. Then, F is a bounded linear surjection.*

We'd like to find an orthonormal set for which F is also one-to-one, but this will be useful for us nonetheless.

PROOF First, why is F linear? Take $\lambda \in \mathbb{F}$ and $x, y \in H$. Then,

$$\begin{aligned} F(\lambda x + y) &= \{(\lambda x + y)_\alpha\}_{\alpha \in \mathcal{J}} \\ &= \{(\lambda x + y, u_\alpha)\}_{\alpha \in \mathcal{J}}. \end{aligned}$$

By the definition of addition and scalar multiplication of functions,

$$\begin{aligned} &= \lambda \{(x, u_\alpha)\}_{\alpha \in \mathcal{J}} + \{(y, u_\alpha)\}_{\alpha \in \mathcal{J}} \\ &= \lambda F(x) + F(y). \end{aligned}$$

Then, Bessel's inequality tells us that

$$\|F(x)\|_{\ell^2(\mathcal{J})}^2 = \sum_{\alpha \in \mathcal{J}} |x_\alpha|^2 \leq \|x\|_H^2,$$

so $\|F\| \leq 1$, and in particular F is bounded.

All the interesting content in the proof is the surjectivity: let $f \in \ell^2(\mathcal{J})$ and $n \in \mathbb{N} \setminus \{0\}$. We want to make things finite, where we can get a handle on them, so let $\mathcal{J}_n = \{\alpha \in \mathcal{J} : |f(\alpha)| \geq 1/n\}$. Then, using $|\cdot|$ to denote cardinality of a set,

$$|\mathcal{J}_n| = \sum_{\alpha \in \mathcal{J}_n} 1 < \sum_{\alpha \in \mathcal{J}_n} (n|f(\alpha)|)^2 \leq n^2 \|f\|_{\ell^2(\mathcal{J})}^2,$$

and the rightmost quantity is finite, so each \mathcal{J}_n is a finite set. Thus, $\mathcal{J} = \bigcup_{i=1}^n \mathcal{J}_n$ is countable, and if $\beta \notin \mathcal{J}$, then $f(\beta) = 0$. Define an $x_n \in H$ by

$$x_n = \sum_{\alpha \in \mathcal{J}_n} f(\alpha) u_\alpha.$$

$\{x_n\}_{n=1}^\infty$ is a Cauchy sequence, because

$$\|x_n - x_m\|^2 = \left\| \sum_{\alpha \in \mathcal{J}_n \setminus \mathcal{J}_m} f(\alpha) u_\alpha \right\|^2 = \sum_{\alpha \in \mathcal{J}_n \setminus \mathcal{J}_m} |f(\alpha)|^2 \leq \sum_{\alpha \in \mathcal{J} \setminus \mathcal{J}_m} |f(\alpha)|^2.$$

Since $\|f\|_{\ell^2(\mathcal{J})}^2$ is finite, then the tail (summing over $\alpha \in \mathcal{J} \setminus \mathcal{J}_n$) goes to 0. Since H is Hilbert, then there's an $x \in H$ such that $x_n \rightarrow x$, and therefore that $F(x_n) \rightarrow F(x)$. Now, we just have to show that $F(x) = f$: for any $\alpha \in I$,

$$\begin{aligned} F(x)(\alpha) &= (x, u_\alpha) = \lim_{n \rightarrow \infty} (x_n, u_\alpha) \\ &= \lim_{n \rightarrow \infty} \sum_{\beta \in \mathcal{J}_n} f(\beta)(u_\beta, u_\alpha) = f(\alpha). \end{aligned} \quad \square$$

Next time, we'll take the maximal orthonormal set, and therefore get a basis, making F one-to-one as well as onto.